

# **Learning from Images**

**Image Features, Descriptors  
and Image Matching**

Master DataScience

Prof. Dr. Kristian Hildebrand  
[khildebrand@bht-berlin.de](mailto:khildebrand@bht-berlin.de)

# Course outline

- **Image descriptors**
  - Global vs. local
  - Harris Corners
  - SIFT
- **Image matching**
  - Distance Metrics
  - FLANN
    - Randomized Kd-Tree
    - K-Means Trees

A photograph of a sunflower field. In the foreground, several sunflowers are visible, their bright yellow petals and dark brown centers contrasting with the green leaves. One prominent sunflower in the center-left is facing towards the camera. The background is filled with more sunflowers and a vast, cloudy sky. The clouds are thick and white, with some darker, more dramatic clouds on the horizon.

What do you see?

**One step back:  
How do we describe an image?  
How do we compare images?**

# Image descriptors

- Descriptor - a piece of stored information that is used to identify an item in an information storage and retrieval system.
  - Descriptors save pertinent information, saving the processing time in future queries
  - Descriptors can save image features that are essential for search and comparison

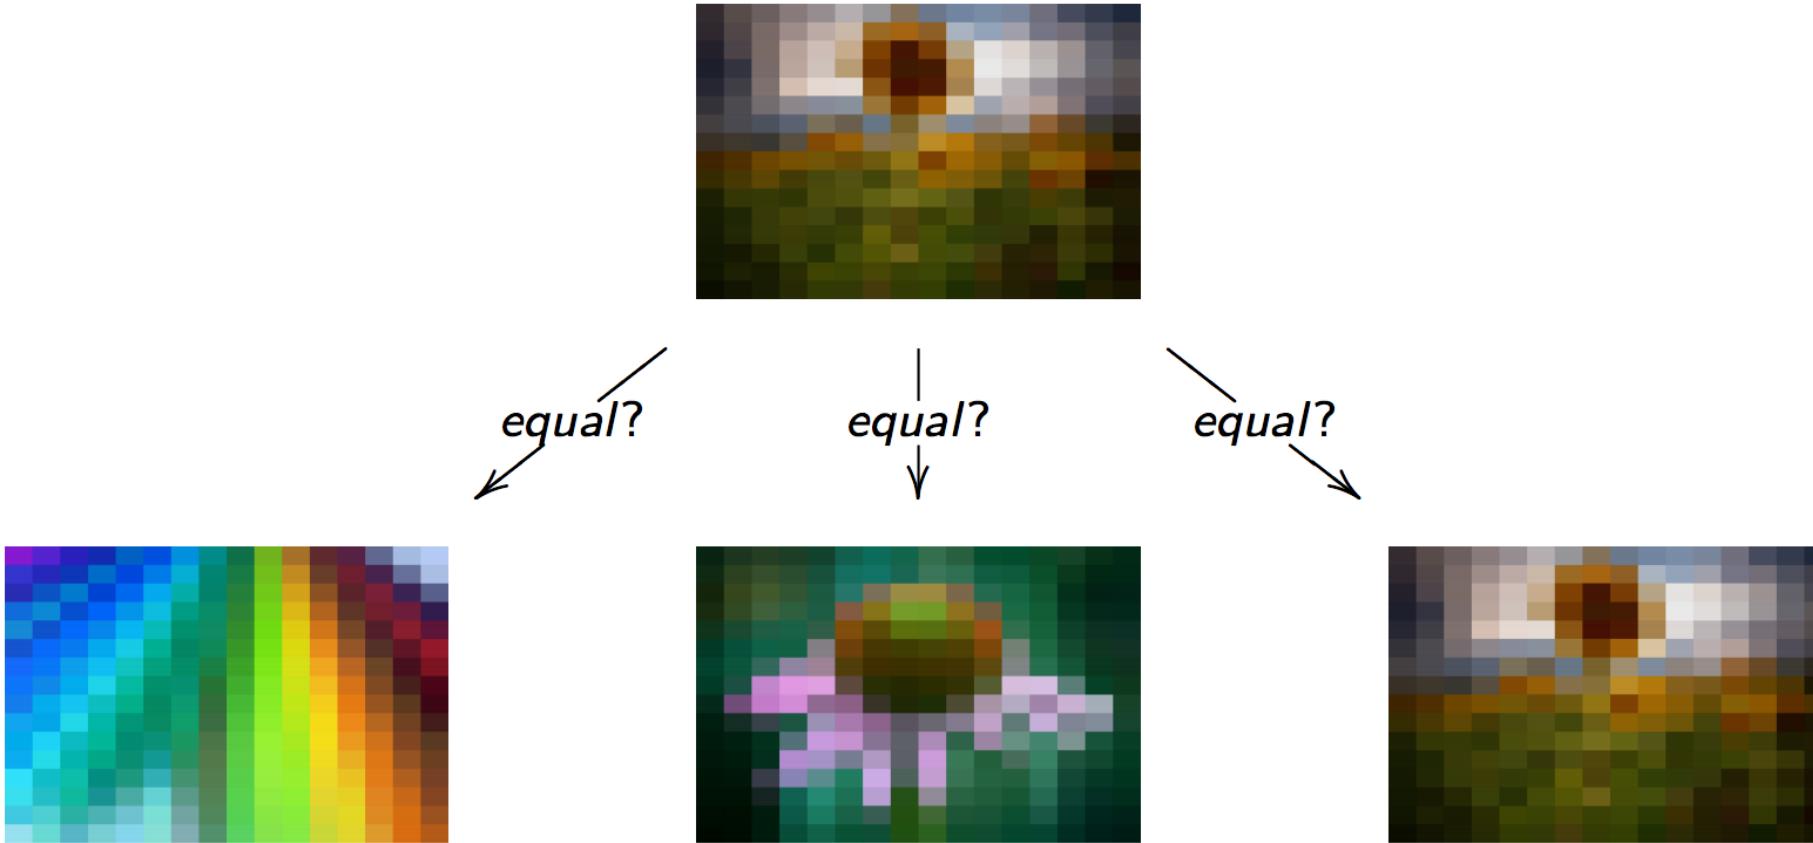
# Global image descriptors



<https://groups.csail.mit.edu/vision/TinyImages/>

<http://www.flickr.com/photos/alphageek/2759567956/>

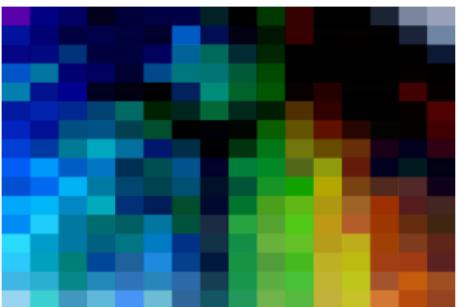
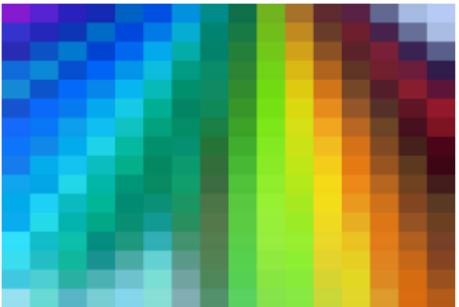
# Global image descriptors



<http://www.flickr.com/photos/40513596@N00/71762740>

<http://www.flickr.com/photos/alphageek/2993813155/>

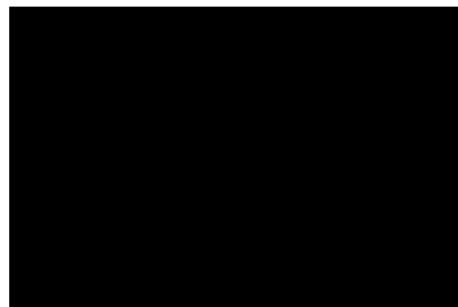
# Comparison example



14862080



437120

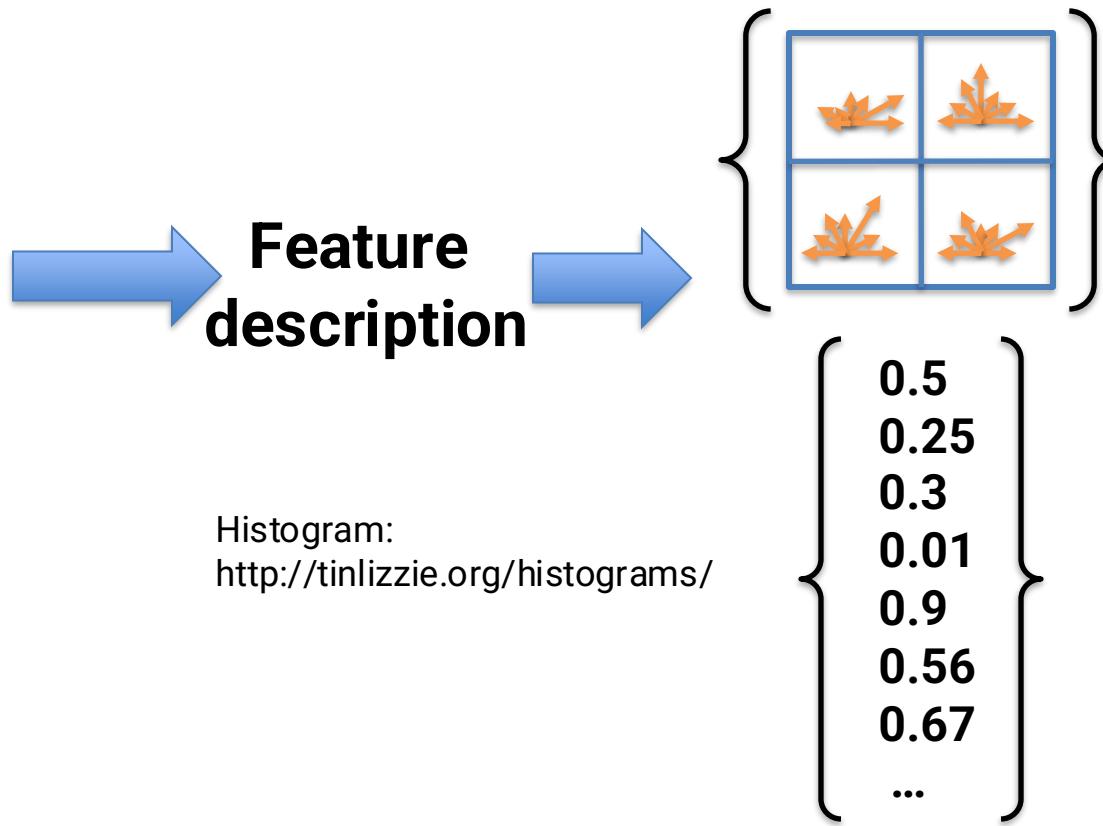


0

# Global vs Local Features

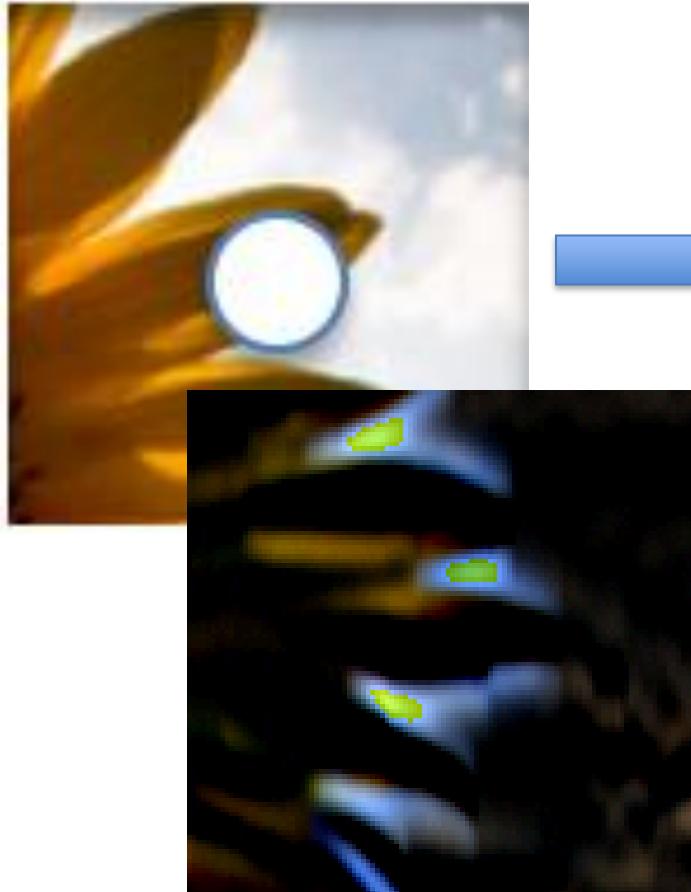


# Image as a vector

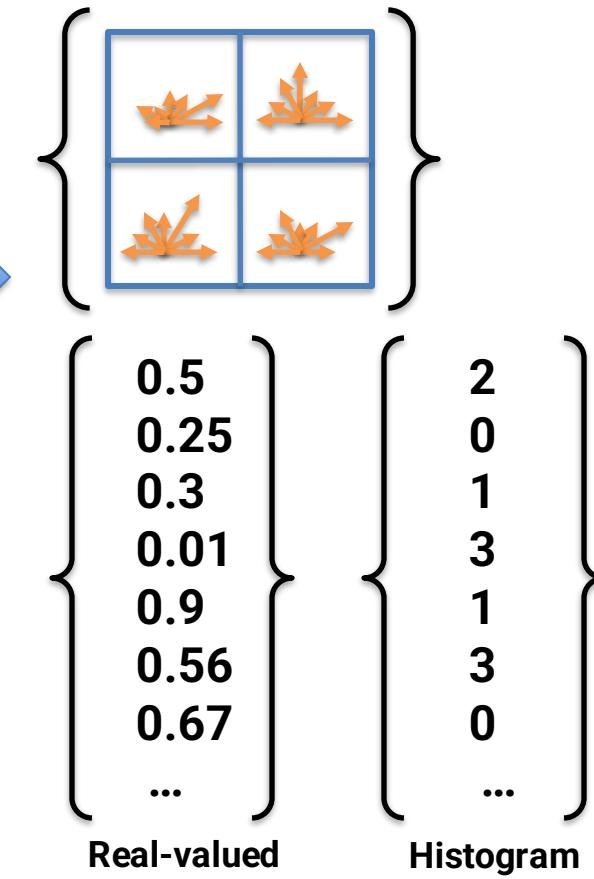


# Image as a vector

What is the advantage of local feature information?



Feature  
description



# Why local features?

- Locality
  - features are local, so robust to occlusion and clutter
- Distinctiveness:
  - can differentiate a large database of objects
- Quantity
  - hundreds or thousands in a single image
- Efficiency
  - real-time performance achievable
- Generality
  - exploit different types of features in different situations

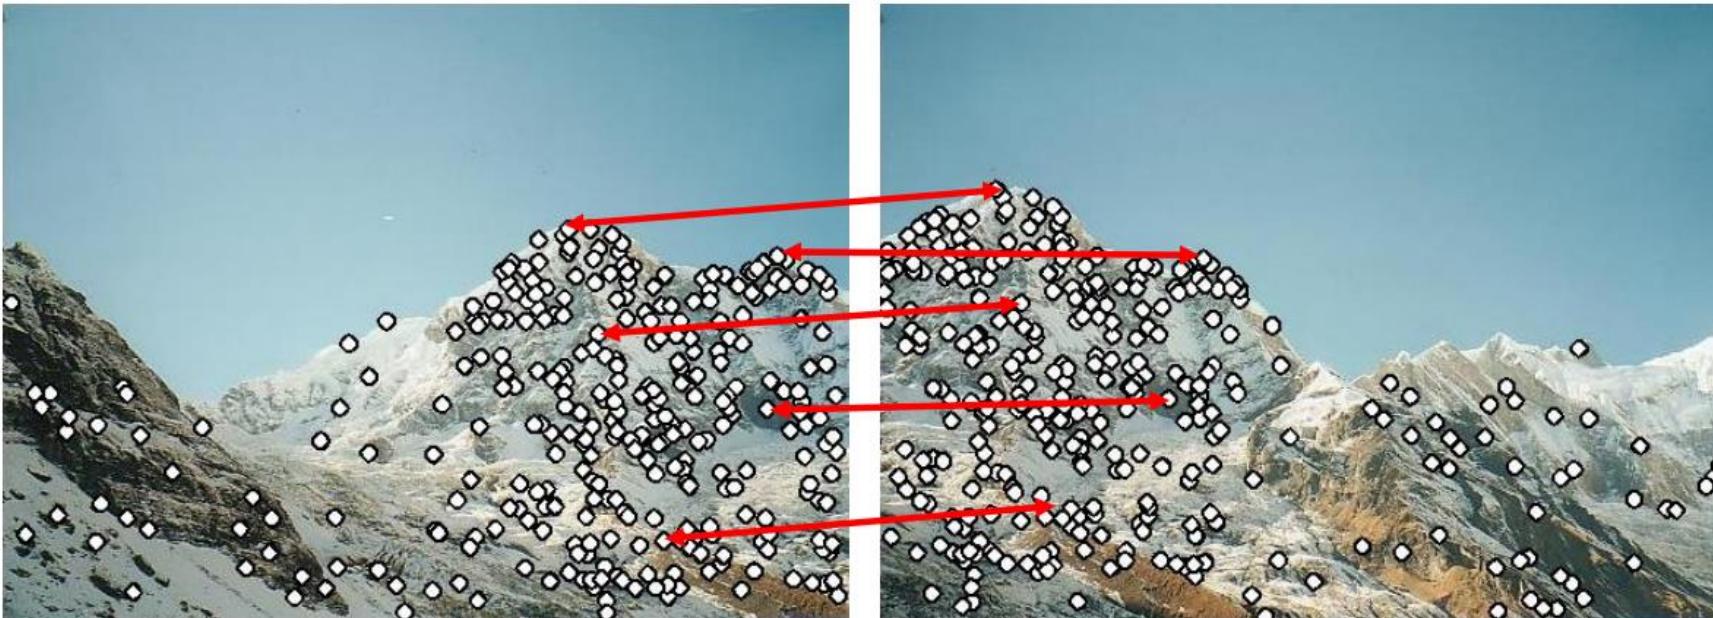
# Applications

- Features are used for:
  - Image alignment (e.g., panoramic mosaics)
  - Object recognition
  - 3D reconstruction (e.g., stereo)
  - Motion tracking
  - Indexing and content-based retrieval
  - Robot navigation
  - ...

**What is a keypoint?  
Or  
Where to extract features?**

# What is a good feature?

Need to detect that feature in the next frame again



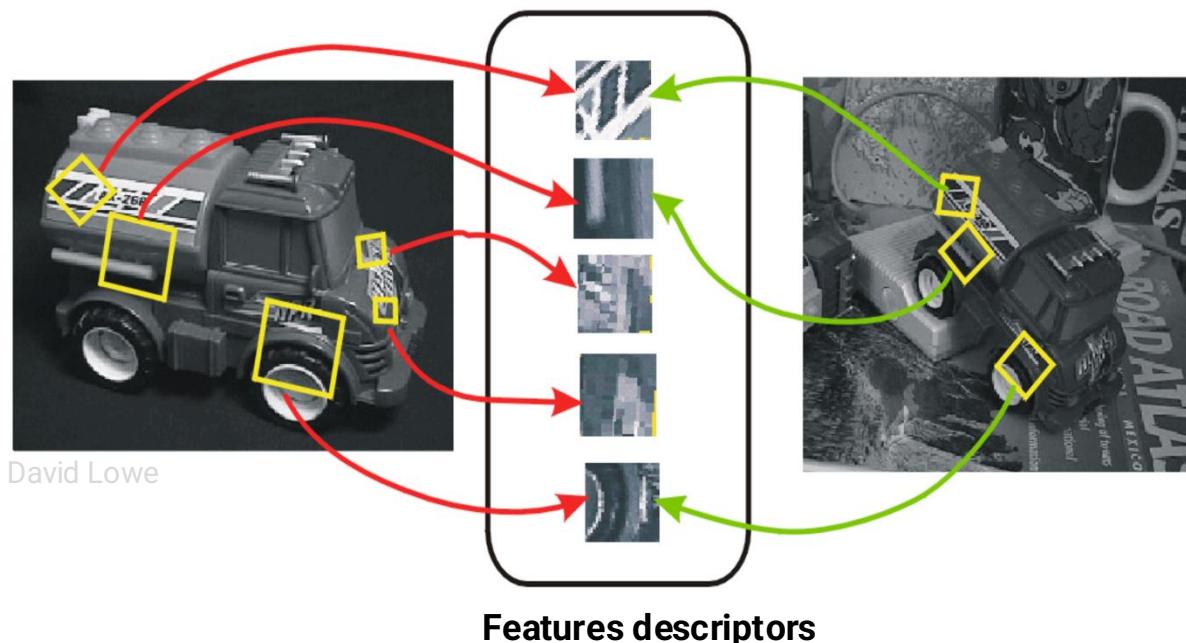
David Lowe

# What makes a good feature?

- Want uniqueness
  - Leads to unambiguous matches in other images
- Look for “interest points”
- Stable under lighting and viewpoint changes
- Keypoints that are similarity- and affine-invariant, or at least stable
- Enable accurate matching of keypoints between images
  - Object recognition / Marker detection / image registration

# Transformation invariance

- Want features that are invariant to transformations (affine-invariant)
  - geometric invariance: translation, rotation, scale
  - photometric invariance: brightness, exposure, ...



# Feature Detection / Location

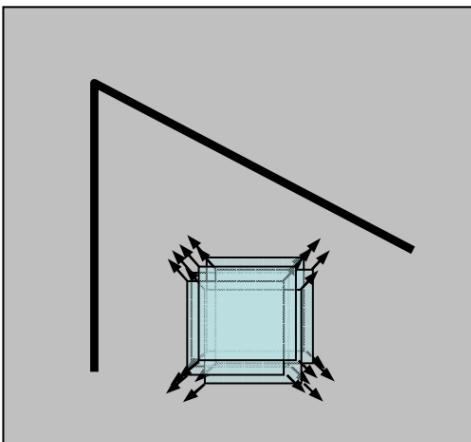
- What could be a possible detector?
  - random?
  - uniform?
  - edges?
  - **corners?**
- corners – intersection of edges
  - point where direction of intersections change
  - gradient of the image (in both directions) has high variation
    - this can be detected = we look for variation

# **Harris Corner Detector**

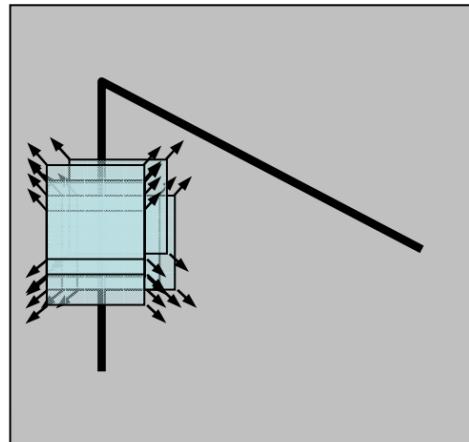
# Harris corner detector – The basic idea

- C.Harris, M.Stephens. “A Combined Corner and Edge Detector”. 1988
- We should easily recognize the point by looking through a small window
- Shifting a window in any direction should give a large change in intensity

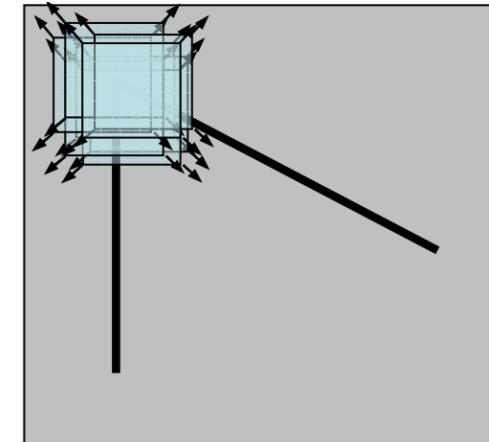
<https://courses.cs.washington.edu/>



“flat” region:  
no change in all directions



“edge”: no change  
along the edge direction



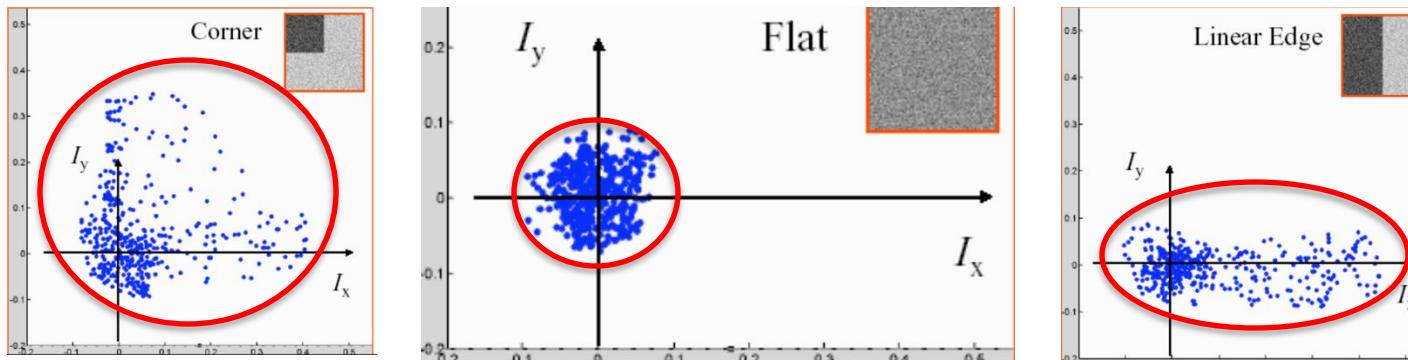
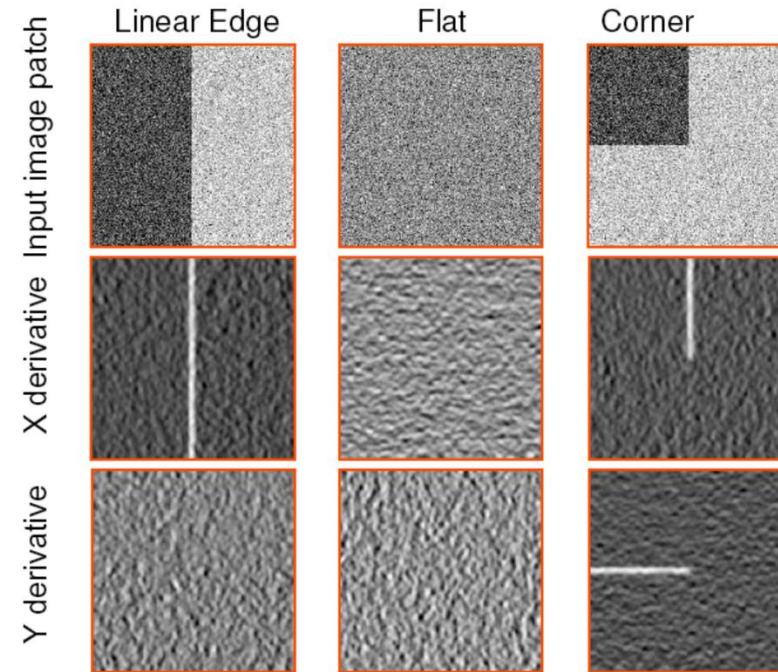
“corner”: significant  
change in all directions

# **Harris corner detector**

## **Demo**

# Harris corner detector – The intuitive way

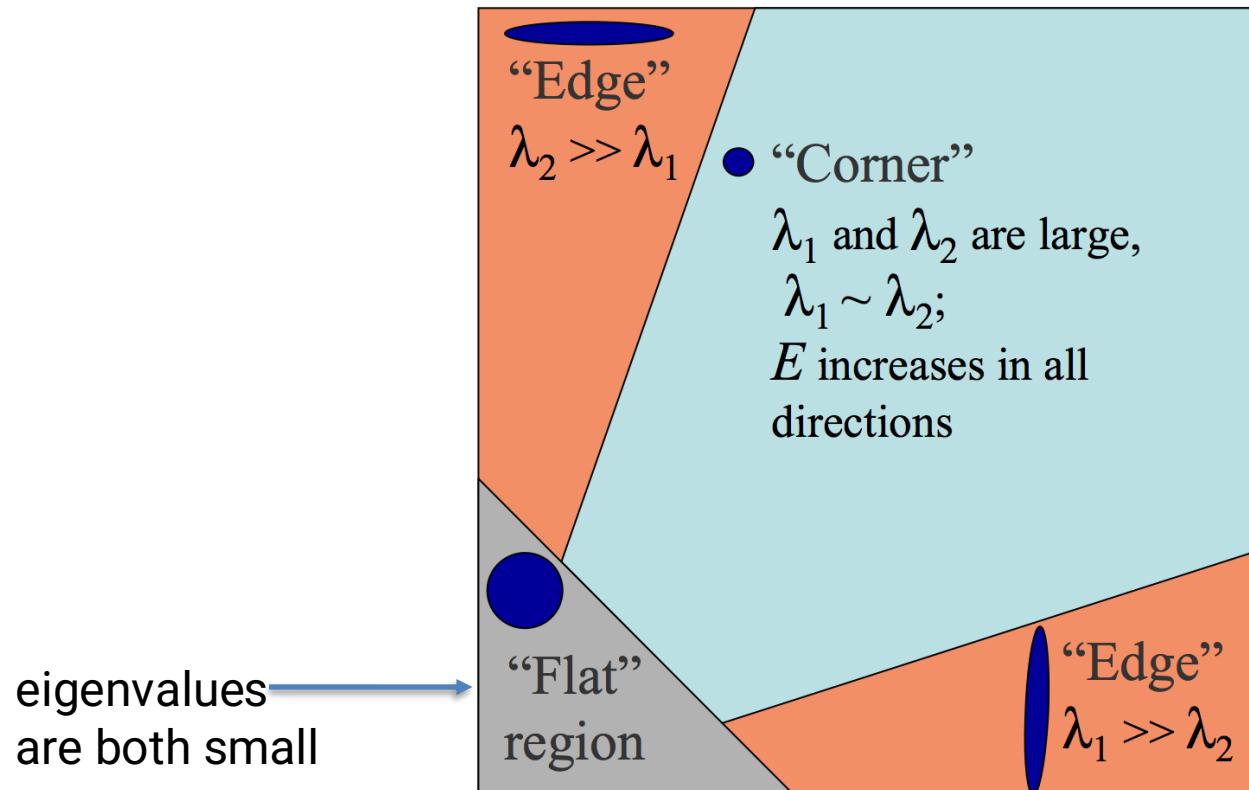
- Treat gradient vectors as a set of  $(dx, dy)$  points with a center of mass defined as being at  $(0,0)$
- Fit an ellipse to that set of points via scatter matrix
- Analyze ellipse parameters for varying cases



# Harris corner detector

- Classification of image points using **eigenvalues of M**

???

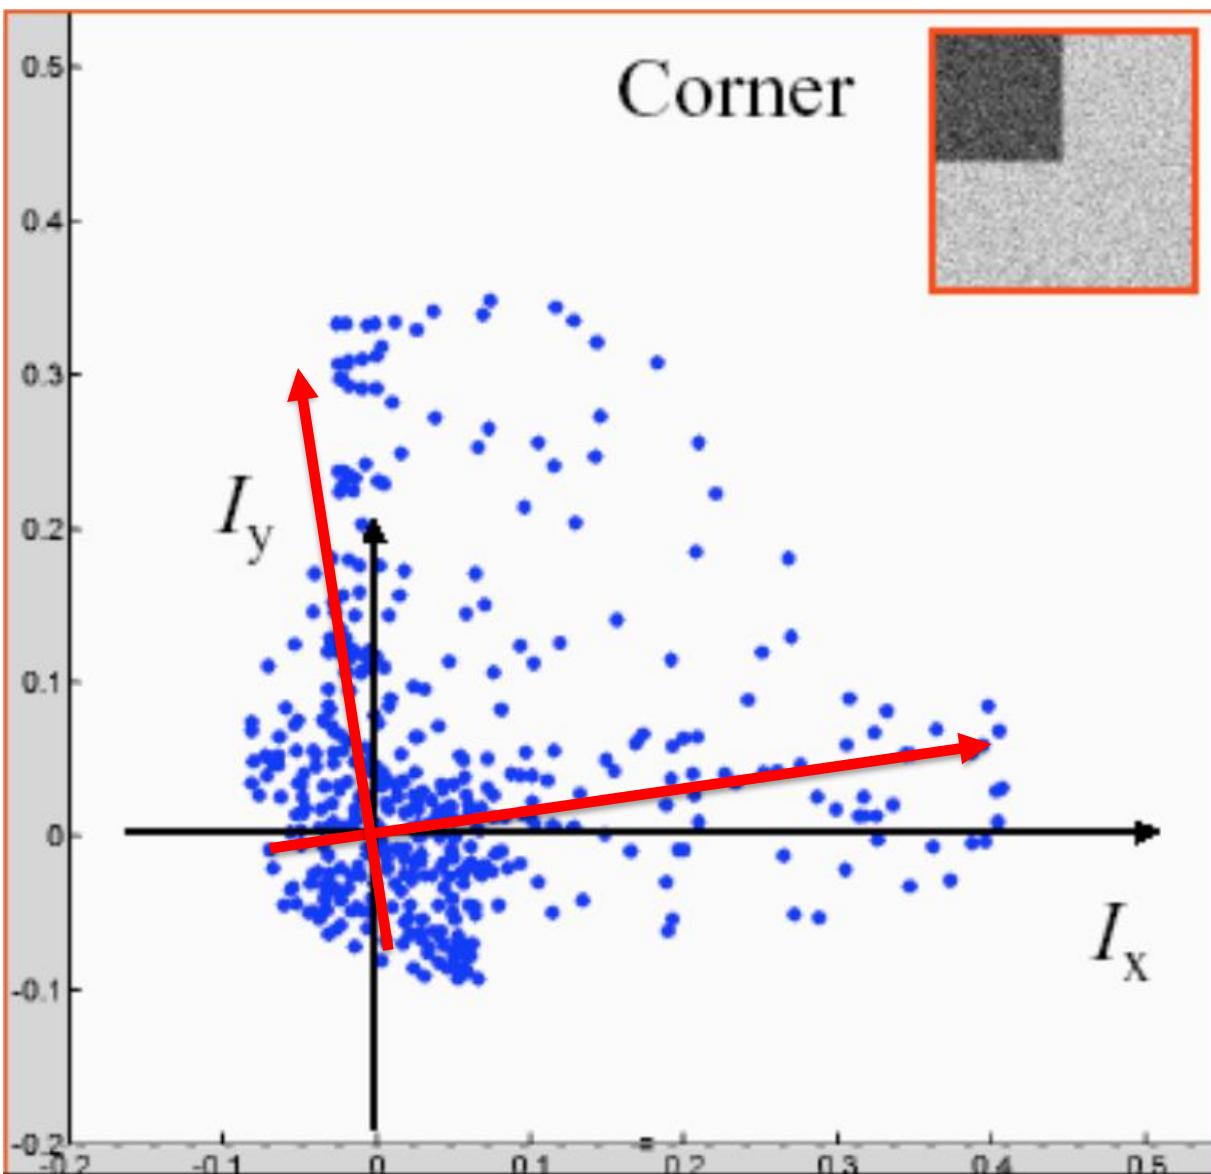


<https://courses.cs.washington.edu/>

**EigenWhat?**

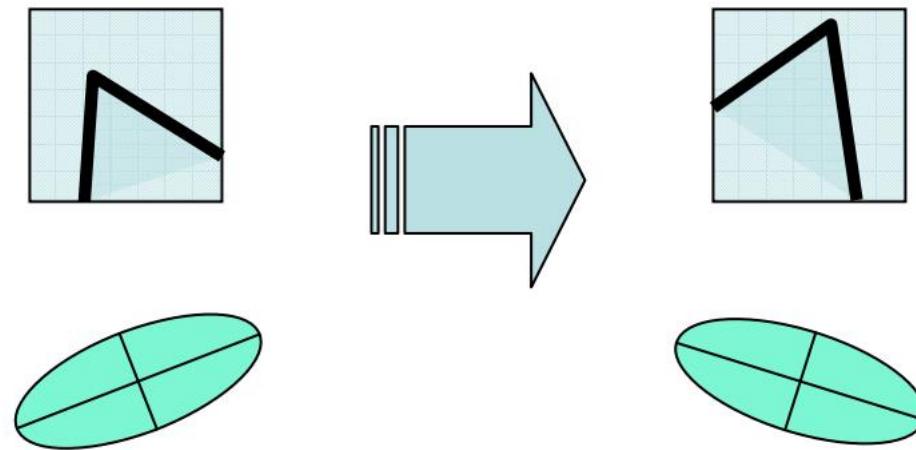
**Eigenvectors / Eigenvalues**

# Eigenvectors - Definition



# Harris corner detector – Rotation invariance

- Ellipse rotates but its shape (i.e. eigenvalues) remains the same



<https://courses.cs.washington.edu/>

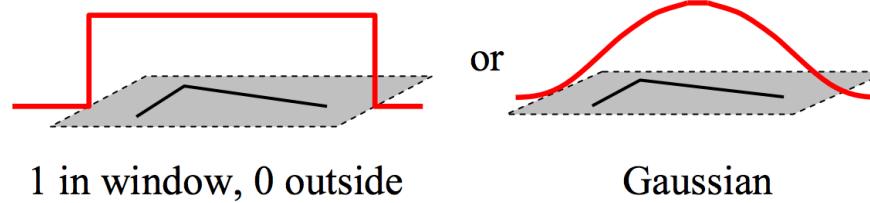
# Harris corner detector – The math

- detect change of intensity for a window shift

$$E(u, v) = \sum_{x,y}^w [I(x + u, y + v) - I(x, y)]^2$$

Window function      Shifted intensity      Intensity

- Window function  $w(x,y)$ :



<https://courses.cs.washington.edu/>

- Maximize variation within window

# Harris Corner Detector – the Math

First order approximation from Taylor Series

$$f(x + u, y + v) \approx f(x, y) + uf_x(x, y) + vf_y(x, y)$$

$$\begin{aligned} E(u, v) &= \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2 \\ &\approx \sum_{x,y} w(x, y) [I(x, y) + uI_x(x, y) + vI_y(x, y) - I(x, y)]^2 \\ &= \sum_{x,y} w(x, y) u^2 I_x^2 + 2uv I_x I_y + v^2 I_y^2 \\ &= \sum_{x,y} (u \quad v) \begin{pmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{pmatrix} \begin{pmatrix} u \\ v \end{pmatrix} \end{aligned}$$

rewrite as matrix equation

$$= \begin{pmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{pmatrix}$$

just the product of components of the gradient  $I_x, I_y$   
also known as structure tensor -> Matrix M

# Harris corner detector – Efficient implementation

- One way to compute the eigenvalues efficiently

$$\lambda_{min} \approx \frac{\lambda_1 \lambda_2}{\lambda_1 + \lambda_2} = \frac{det(M)}{trace(M)}$$

- To avoid computing the eigenvalues which is computationally expensive one uses an approximation to determine if point is corner:

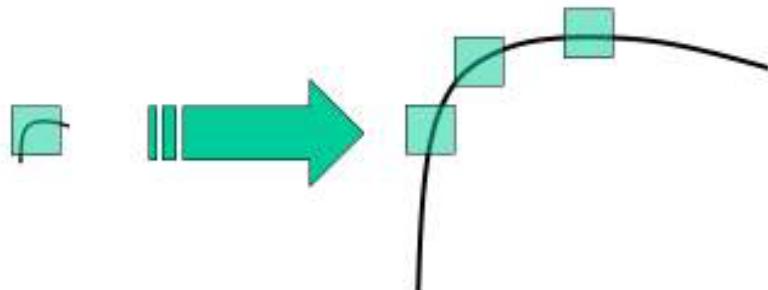
$$R = det(M) - k(trace(M))^2$$

Implementation is done as homework

# **Scale invariant feature transform (SIFT)**

# Scale invariant feature transform (SIFT)

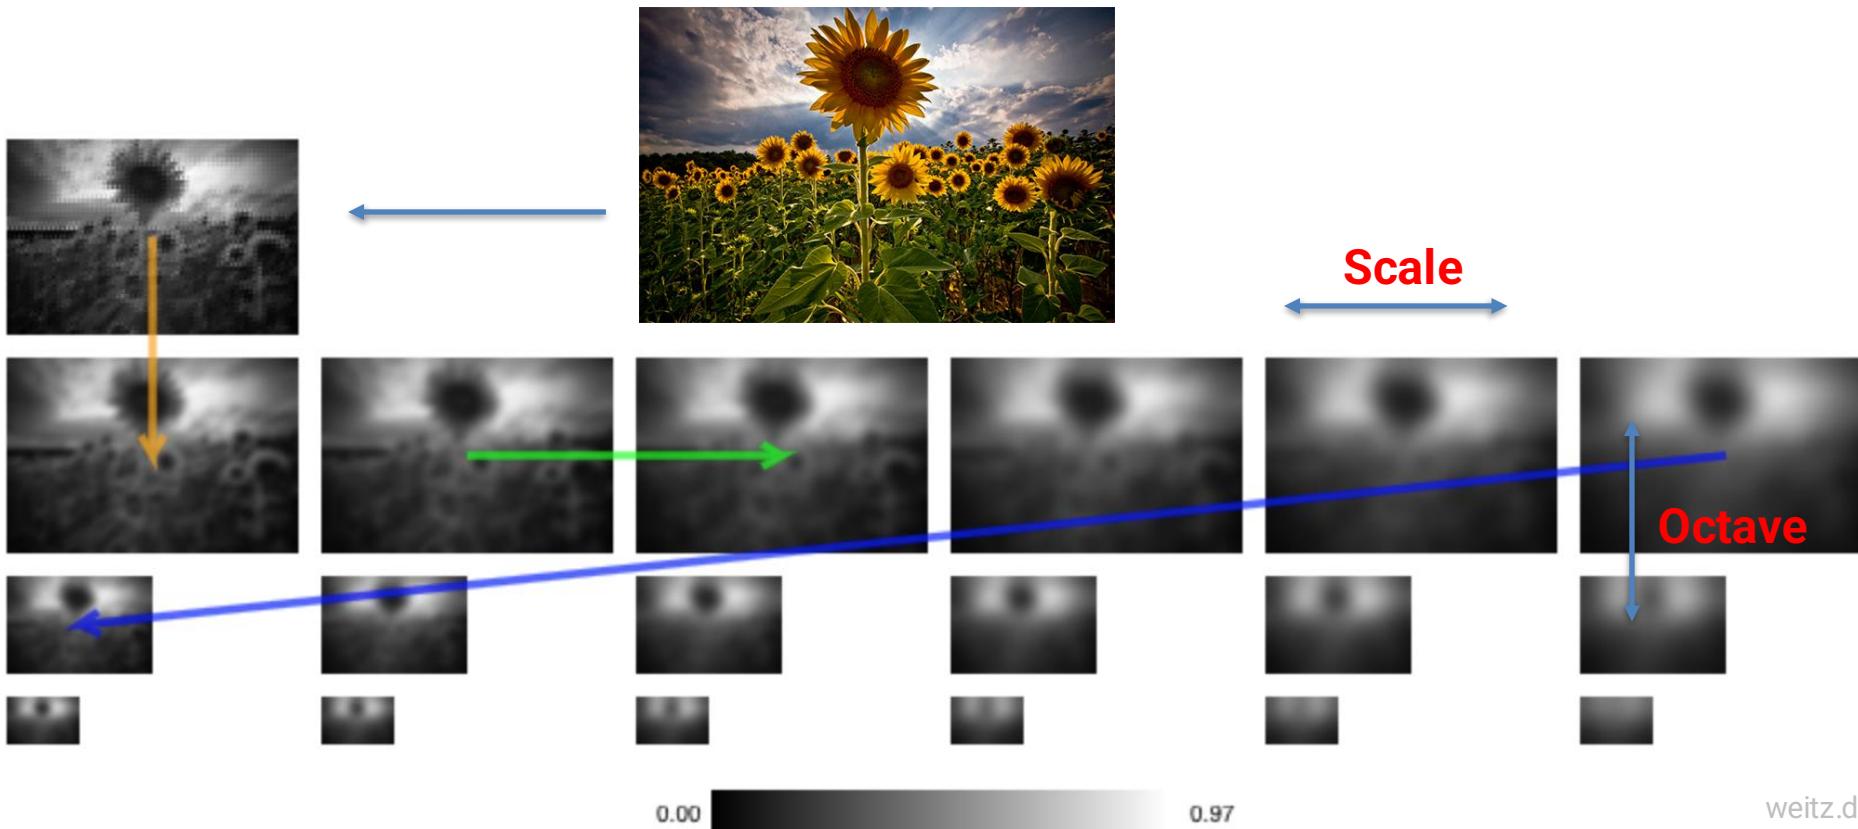
- Distinctive Image Features from Scale-Invariant Keypoints, David Lowe, 2004
- corner detectors are rotation-invariant
  - even if the image is rotated, one can find the same corners
  - corners remain corners in rotated image
- this is not true for scaling



- Harris detector is not scale-invariant

opencv.org

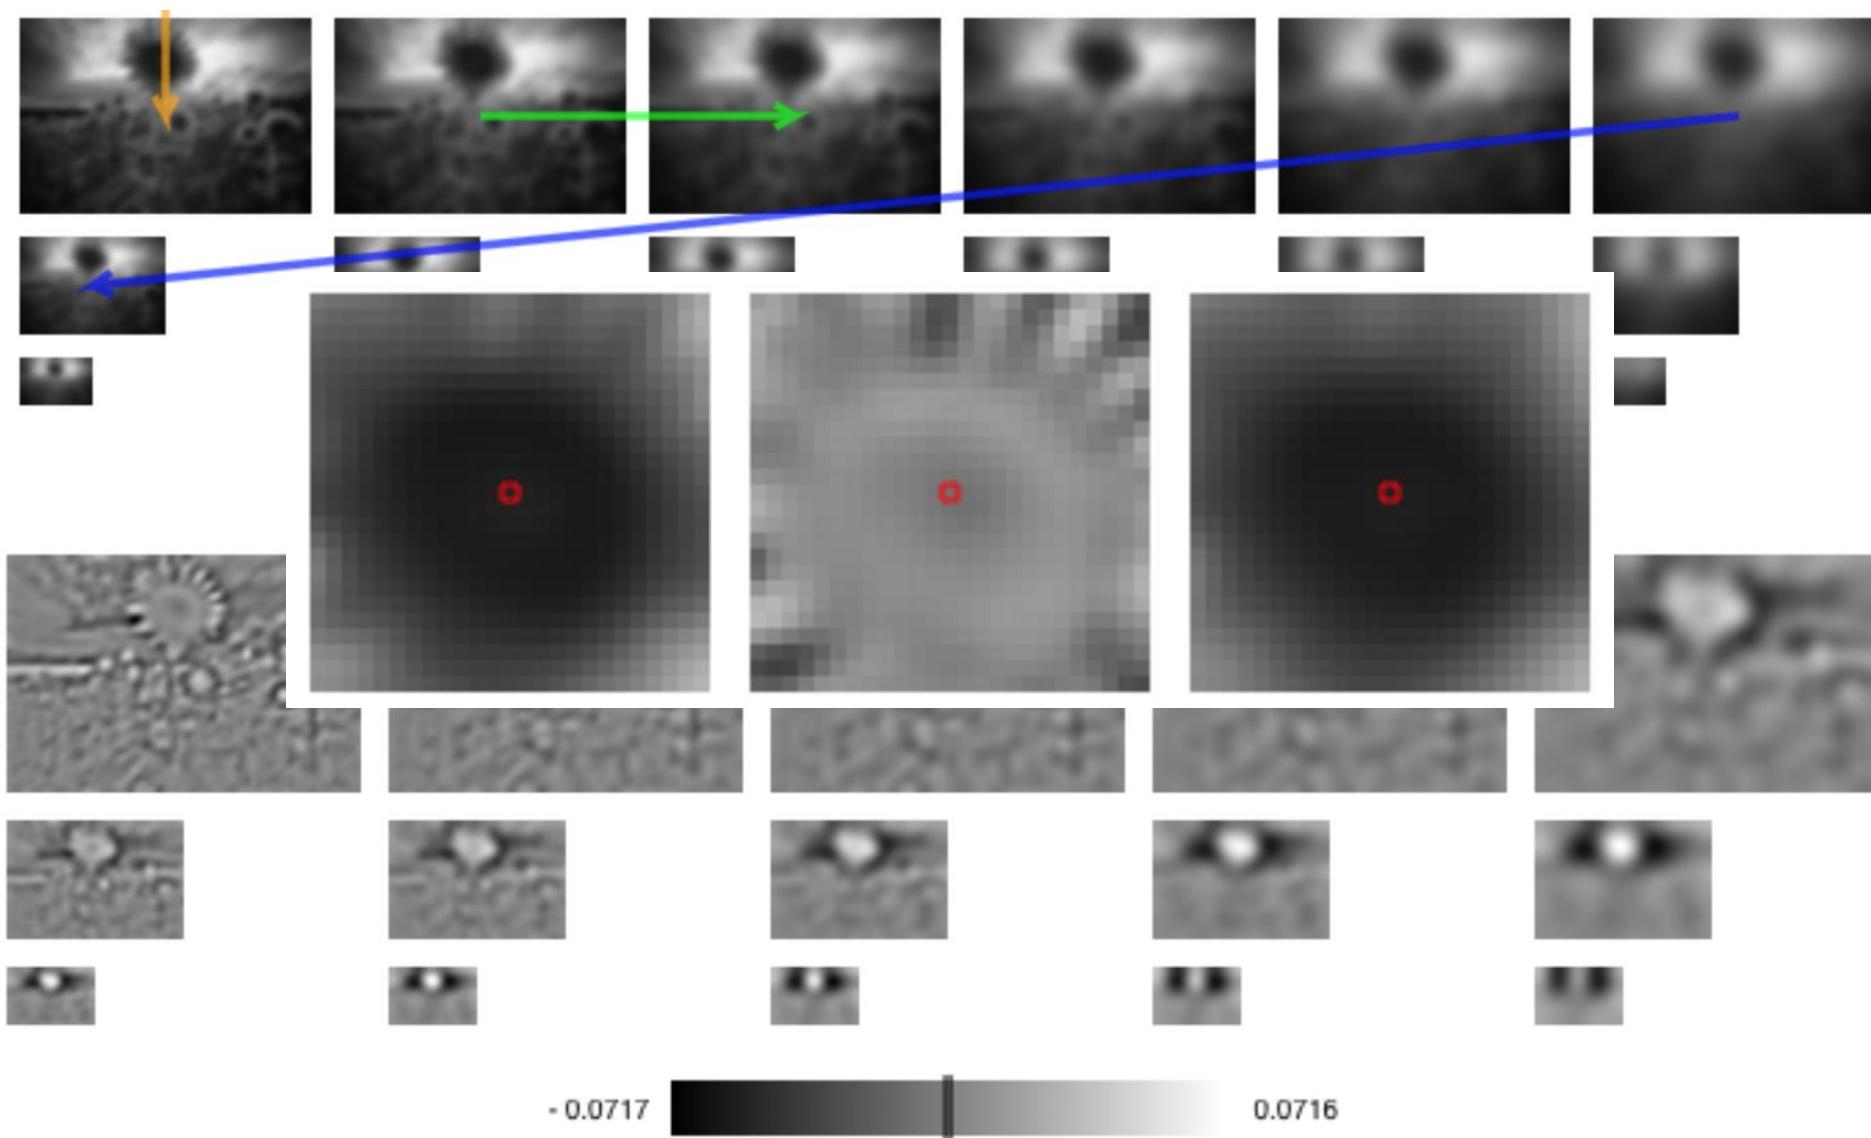
# SIFT - Scale-space Extrema Detection



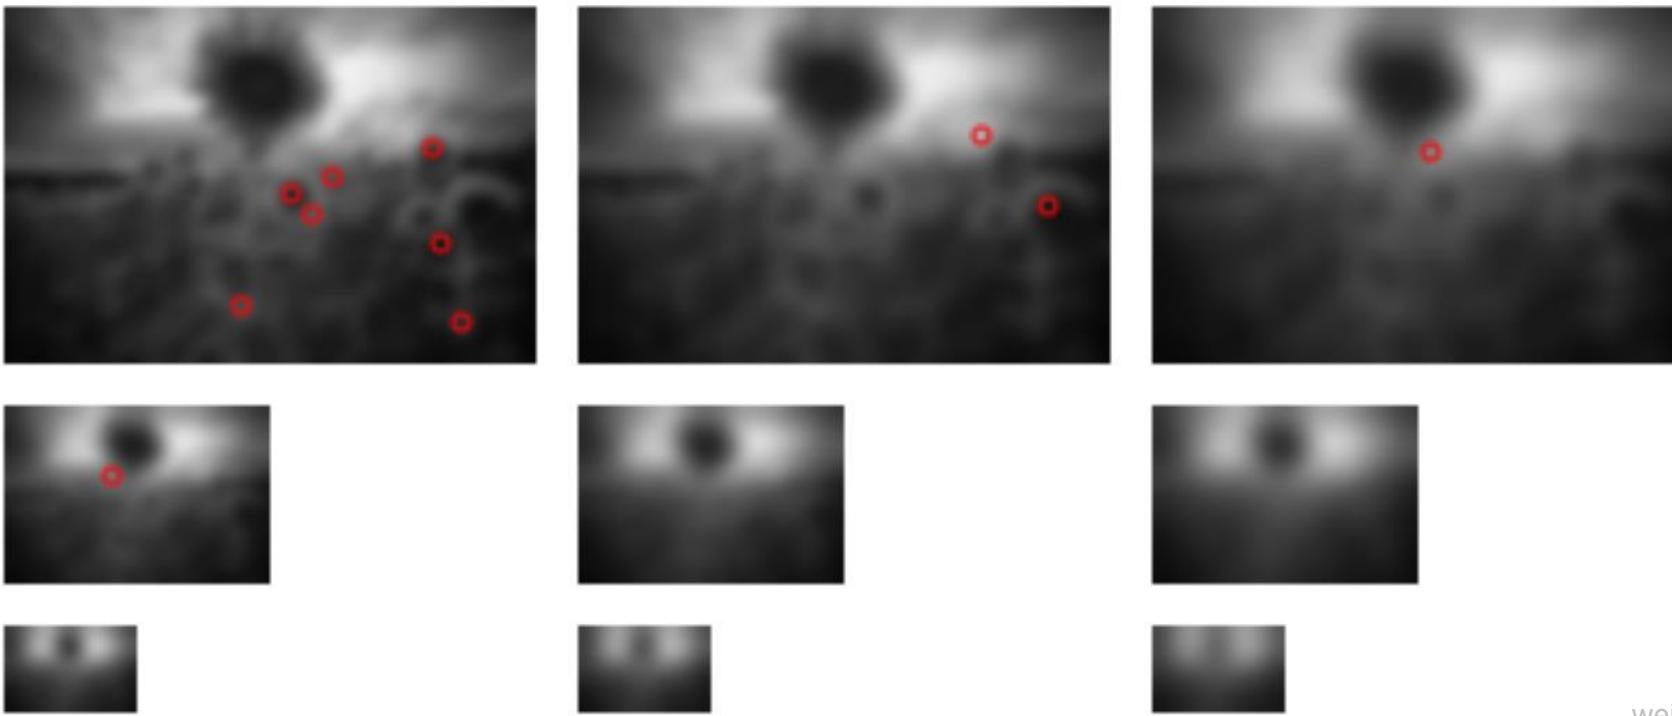
**Scale-space using bilinear interpolation  
and gaussian convolution**

$$G(x, y, \sigma) = \frac{1}{2\pi\sigma^2} e^{-(x^2+y^2)/2\sigma^2}$$

# SIFT - Difference of Gaussians



# SIFT - Keypoints localization



weitz.de

- one pixel in an image is compared with its 8 neighbours as well as 9 pixels in next scale and 9 pixels in previous scales
- If it is a local extrema, it is a potential keypoint
- It basically means that keypoint is best represented in that scale

# **Feature**

## **Description of the image window around the keypoint**

# SIFT - Keypoint Descriptor

- 16x16 neighbourhood around the keypoint
- divided into 16 sub-blocks of 4x4 size
- each sub-block, 8 bin orientation histogram is created
  - total of 128 bin values are available
- represented as a vector to form keypoint descriptor

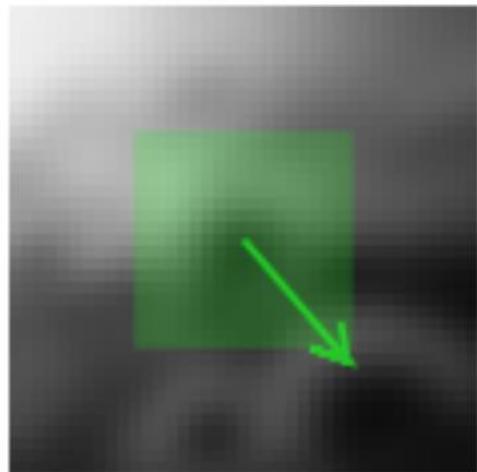
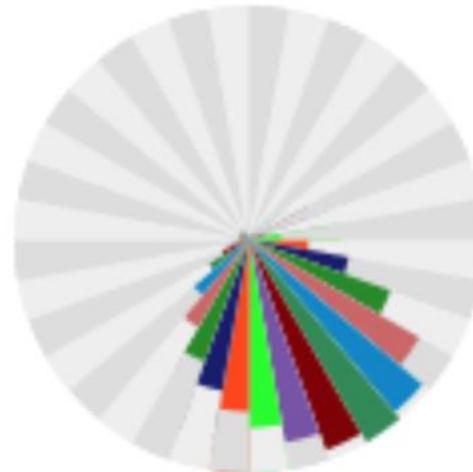


Image Gradients



Keypoint descriptor



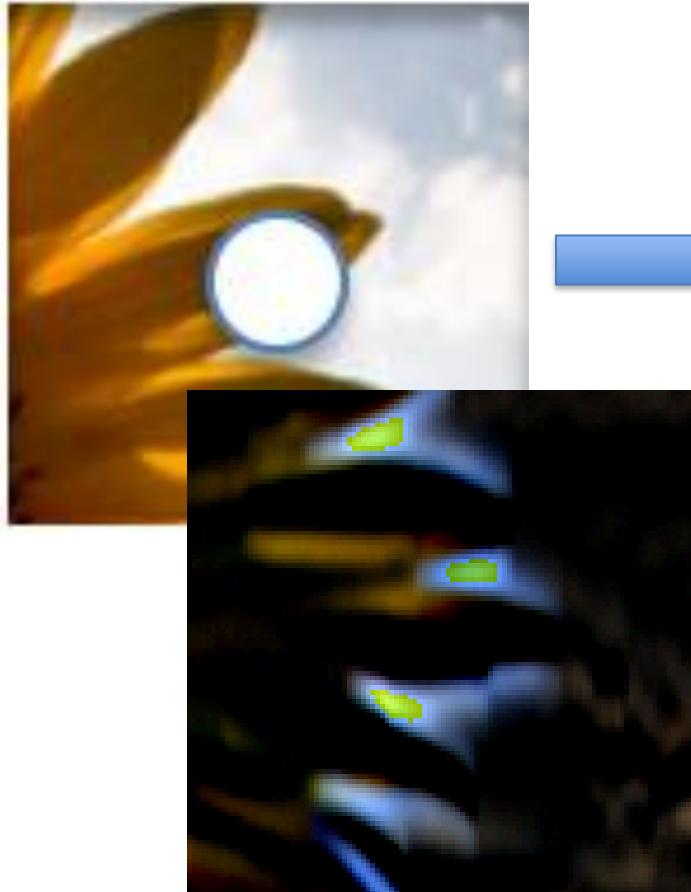
# More feature descriptors and keypoint detectors

- FAST - Features from Accelerated Segment Test
  - Detects corners using the FAST algorithm
- MSER - Maximally stable extremal region extractor.
- ORB - oriented *BRIEF* Keypoint detector and descriptor extractor
- BRISK - Binary Robust Invariant Scalable Keypoints
- SURF -Speeded-Up Robust Features
- FREAK -*Fast Retina Keypoint*

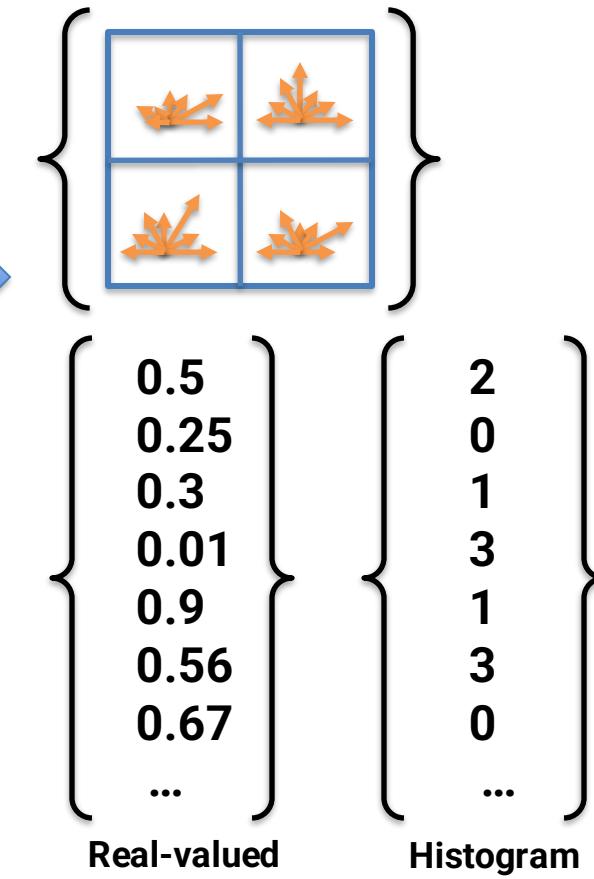
# **Intermediate take away**

# Image as a vector

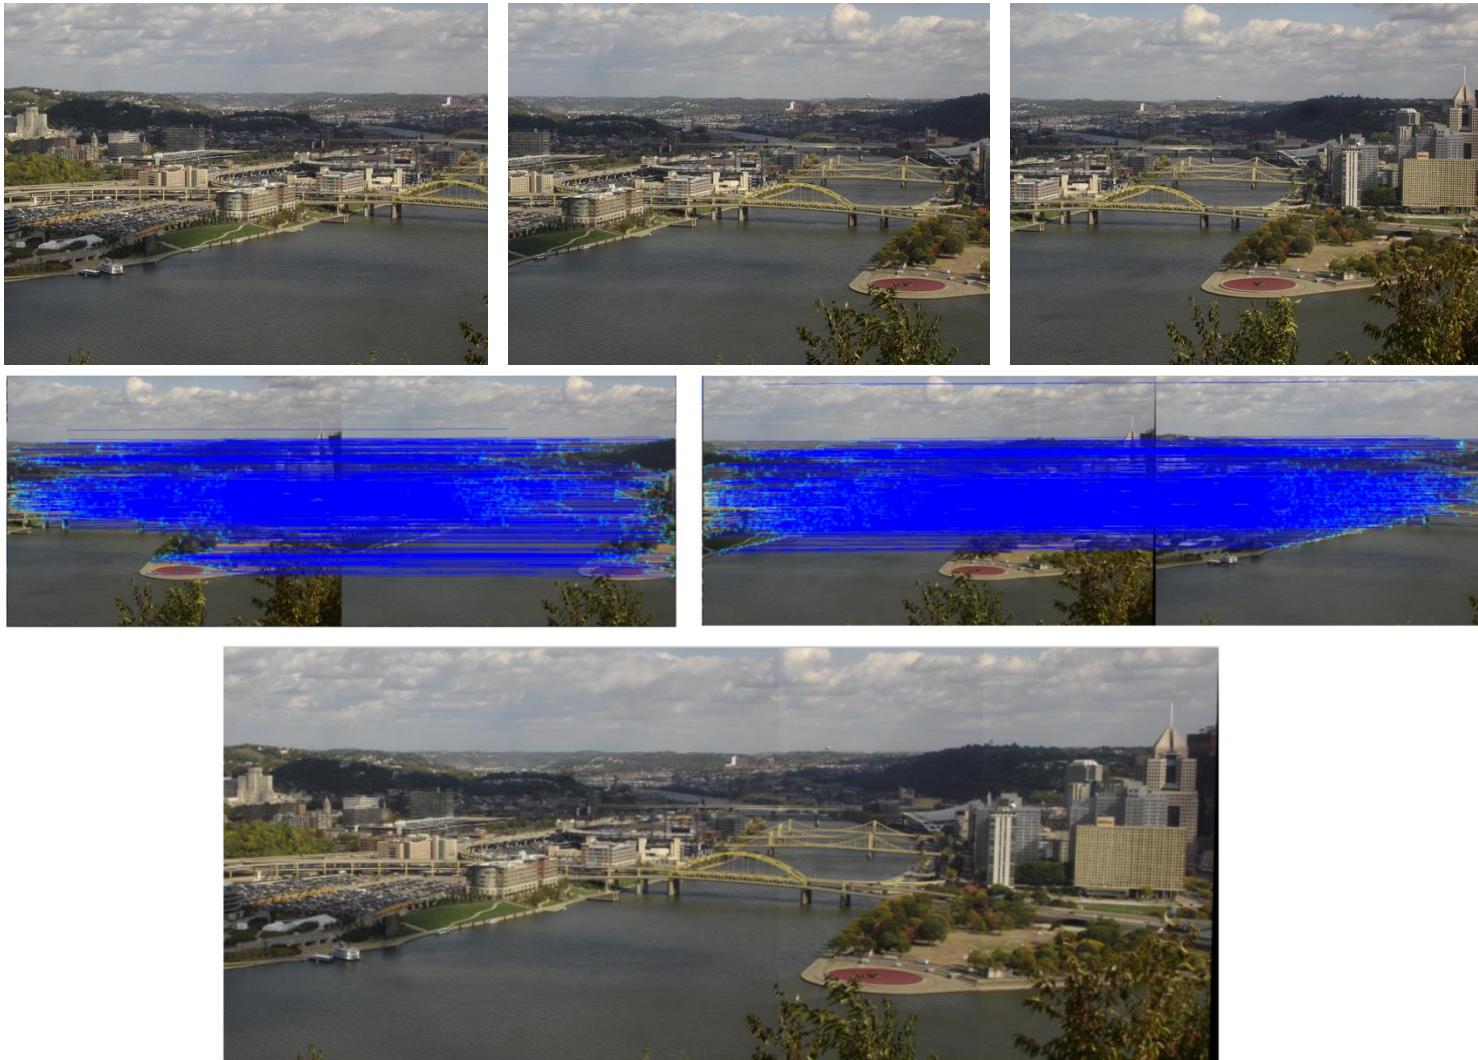
What is the advantage of local feature information?



Feature  
description



# Image feature example applications



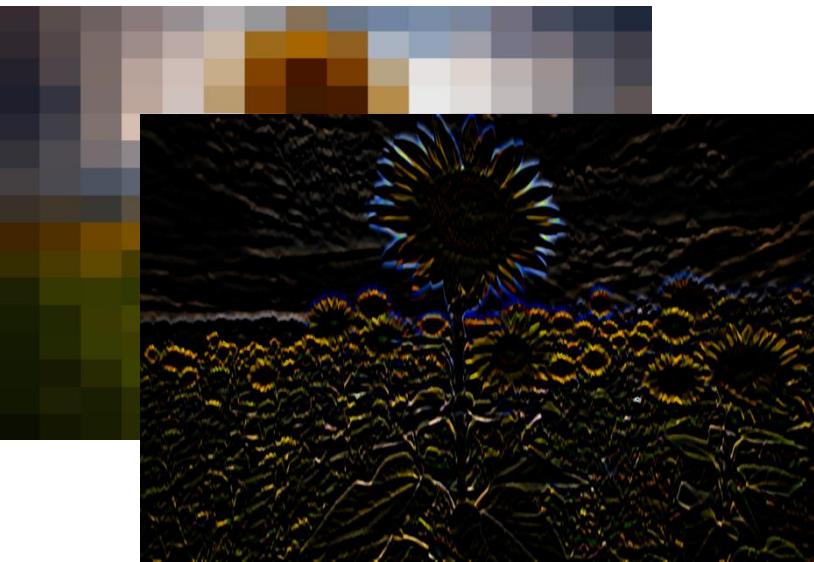
# Image feature example applications



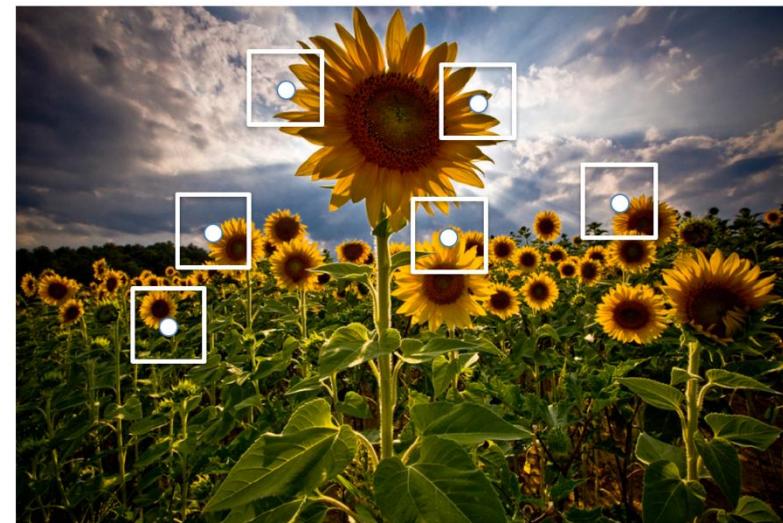
# Image features



Global Descriptor

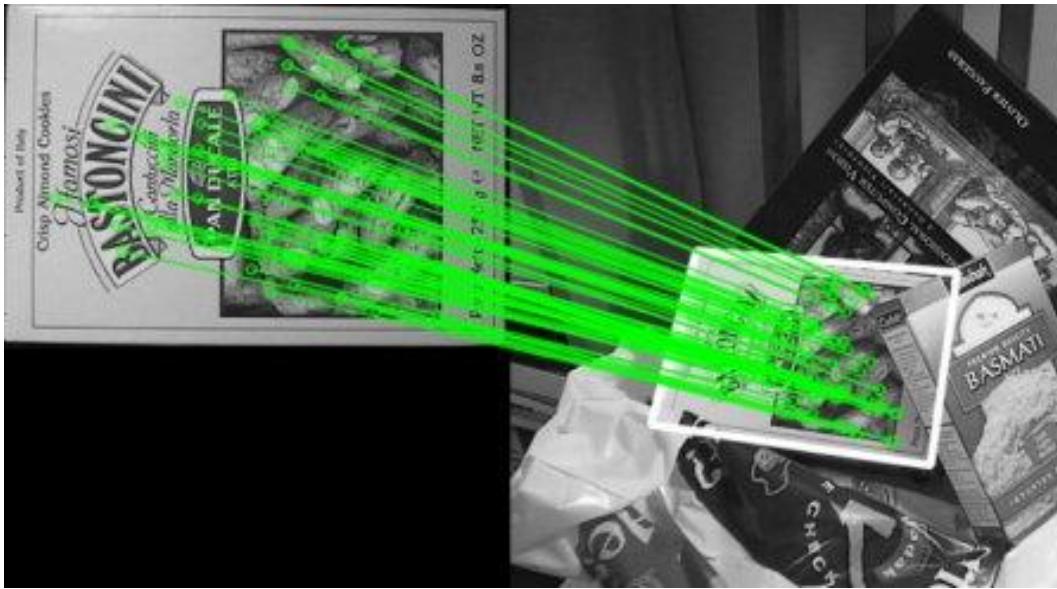


Local Descriptor

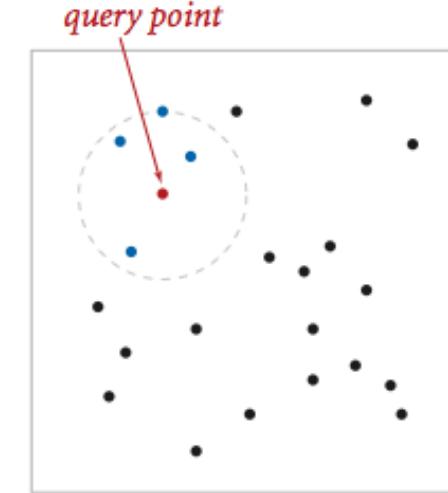
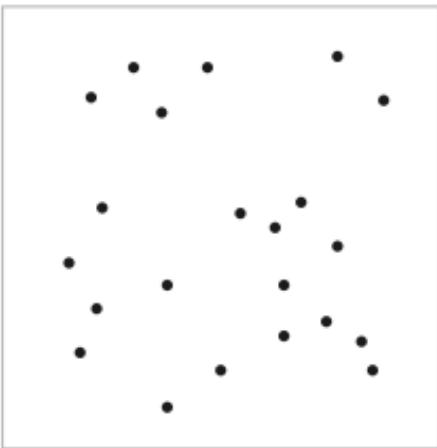
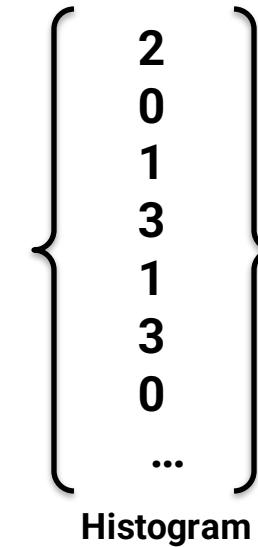


# **Feature Matching**

# Feature Matching



<http://www.ubc.ca>



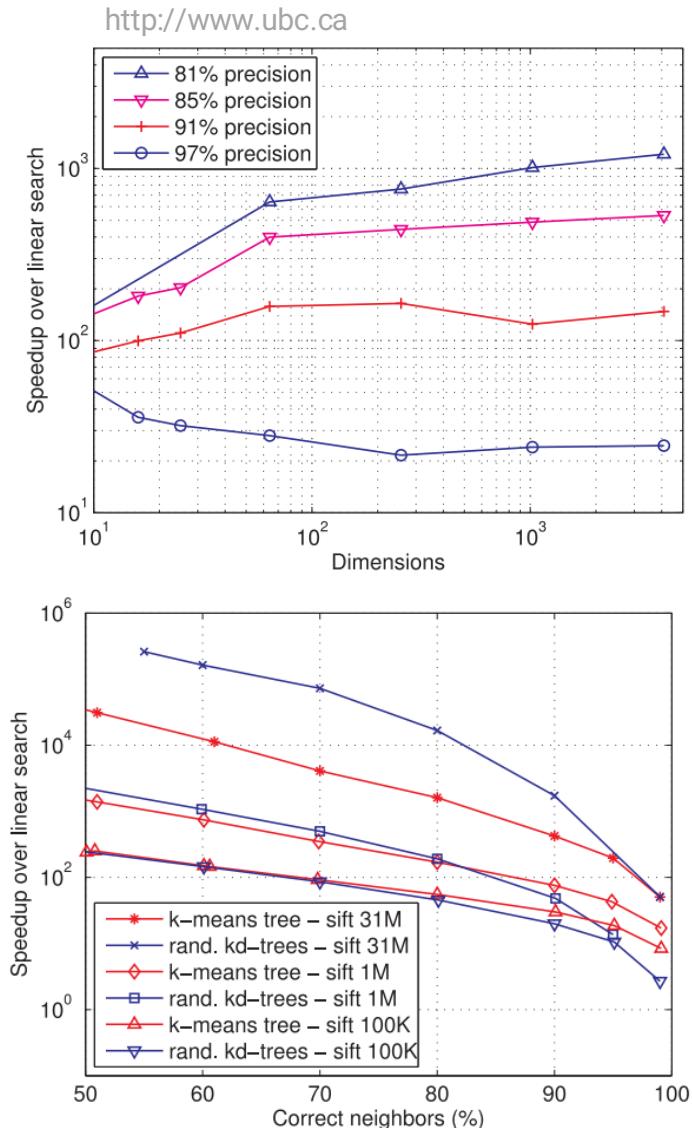
<https://www.cs.princeton.edu>

# Feature Matching

- searching for the **most similar matches to high-dimensional vectors**  
     **nearest neighbor matching**
- **no known exact algorithms** for solving these high-dimensional problems that are faster than linear search  
     **Curse of dimensionality**
- **efficient algorithm** for fast nearest neighbor (NN) matching in large data sets can bring speed improvements of several orders of magnitude
- **FLANN – Fast Library for Approximate Nearest Neighbor**
  - library for performing fast approximate nearest neighbor searches in high dimensional spaces

# Feature Matching

- Obtain **speed improvement by allowing an approximate search**
  - not all the neighbors returned are exact
  - some are approximate but still close to the exact neighbors
- **Approximate nearest neighbor search algorithms provide ~95% correct neighbors**
  - **BUT:** are two or more orders of magnitude faster than linear search



# Feature Matching Applications

- Nearest neighbor search problem is also of major importance for many **applications**:
  - machine learning
  - image- / document-retrieval
  - data compression
  - bio-informatics
  - data analysis
  - computer vision

# Feature Matching

- Given a set of points in metric space  $M$  and a query point  $q$

$$P = \{p_1, p_2, p_3, \dots, p_n\} \quad q \in M$$

- Find elements that is closest to  $q$  with respect to a metric distance  $d$

$$NN(q, P) = \operatorname{argmin}_{x \in P} d(q, x)$$

# Feature Matching

- **How to compare feature vectors in high-dimensional space?**
  - Distance Metric
- How to compare feature vectors in high-dimensional space **efficiently?**
  - Spatial Datastructure
  - Approximate Nearest Neighbor

# **Distance Metric (Vector Space Norm)**

1.  $d \geq 0$
2.  $d(x, y) = d(y, x)$
3.  $d(x, y) \leq d(x, y) + d(y, z)$
4.  $d(x, y) = 0 \rightarrow x = y$

# Distance metric

- **Euclidean distance**
  - ‘straight-line’ distance between points ( $L^2$ -norm)

$$d_{l2}(p, q) = \sqrt{\sum_{i=1}^n (p_i - q_i)^2}$$

- **Manhattan distance**

- L1-norm

$$d_{l1}(p, q) = \sum_{i=1}^n |p_i - q_i|$$

- **Maximum norm**

- $L^\infty$

$$d_{l\infty}(p, q) = \max_i(|p_i - q_i|)$$

# Distance metric

- **Cosine similarity (strictly speaking – this is not a distance metric)**
  - similarity between two vectors (inner product,  $\cos(\text{angle})$  between vectors)
  - resulting similarity ranges from -1 meaning opposite, to 1 meaning the same, with 0 indicating orthogonality
  - in-between values indicate intermediate similarity or dissimilarity
  - distance between histograms

$$a \cdot b = \|a\| \|b\| \cos \theta$$

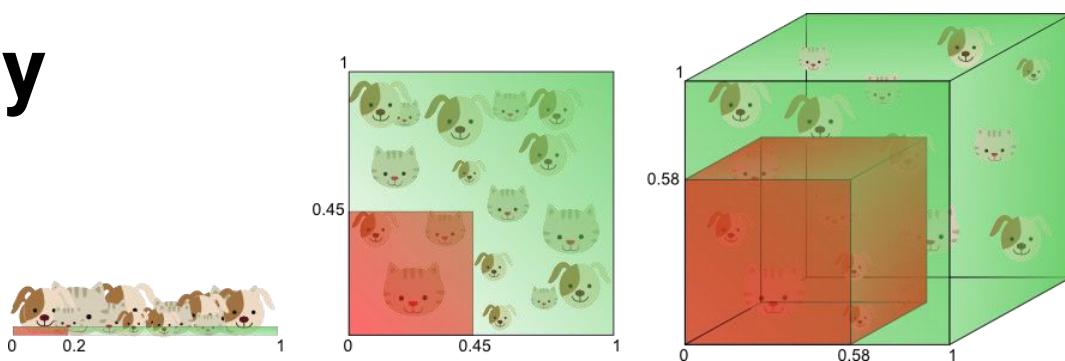
$$\text{similarity} = \cos(\theta) = \frac{a \cdot b}{\|a\| \cdot \|b\|}$$

**Something *not* completely unrelated...**

**Curse of ...**

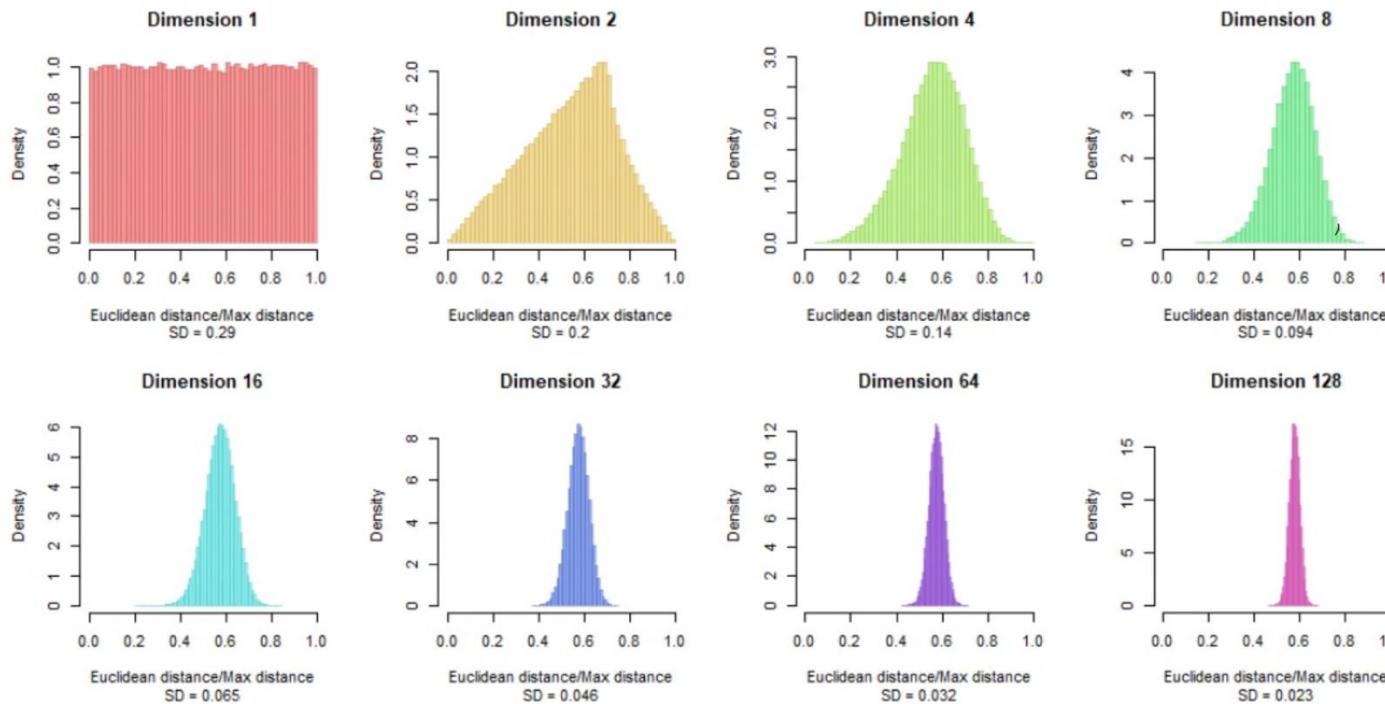


**Dimensionality**



# Curse of dimensionality

- Volume grows rapidly with increase of dimensions
- With increasing dimensions the difference in distances decreases



$$\lim_{d \rightarrow \infty} E \left( \frac{\text{dist}_{\max}(d) - \text{dist}_{\min}(d)}{\text{dist}_{\min}(d)} \right) \rightarrow 0.$$

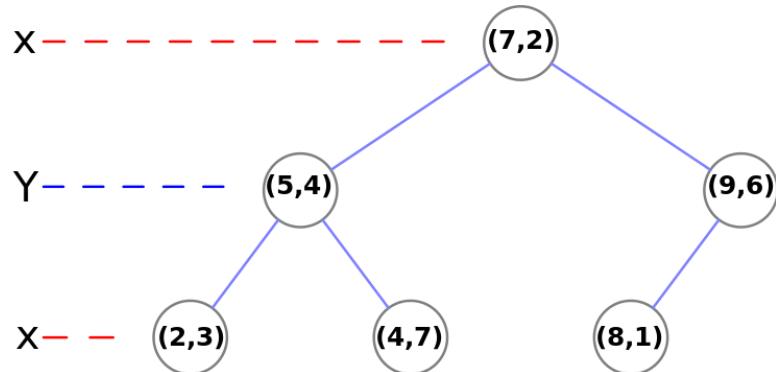
# Feature Matching

- How to compare feature vectors in high-dimensional space?
  - Distance Metric
- **How to compare feature vectors in high-dimensional space efficiently?**
  - Spatial Datastructure
  - Approximate Nearest Neighbor

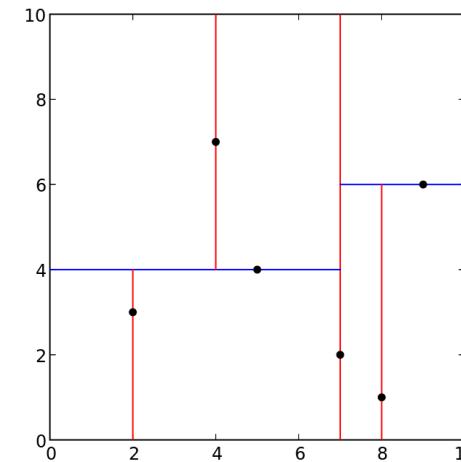
# **Multiple Randomized kd-Trees**

# kd-tree (Recap)

- **spatial data structure** for organizing points in **high dimensionality spaces**
- **binary tree** where each node is a k-d point
- every **node** can be seen **as a plane dividing space in two parts**
- splitting along changing dimensions / axis
- **results in balanced k-d tree**

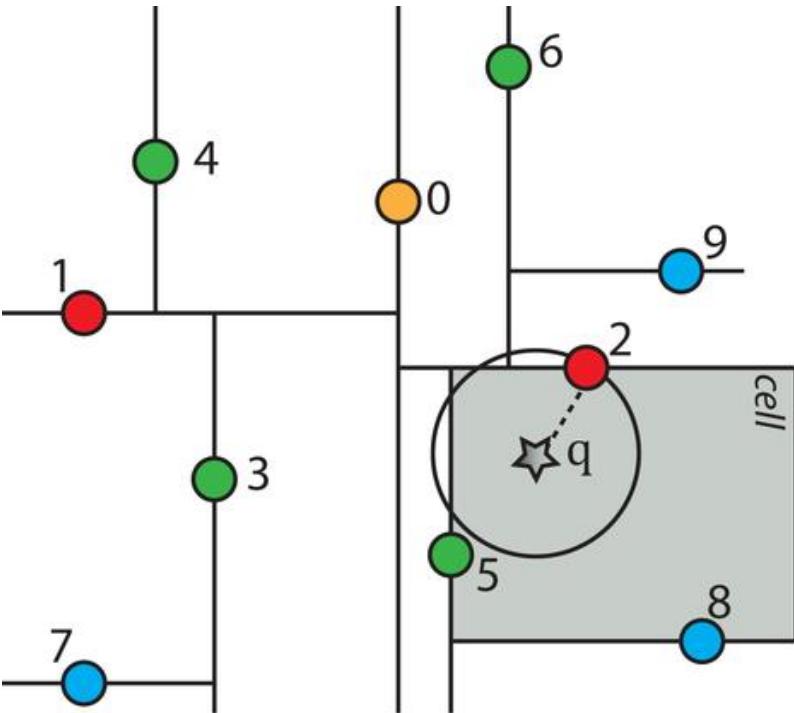


white board example

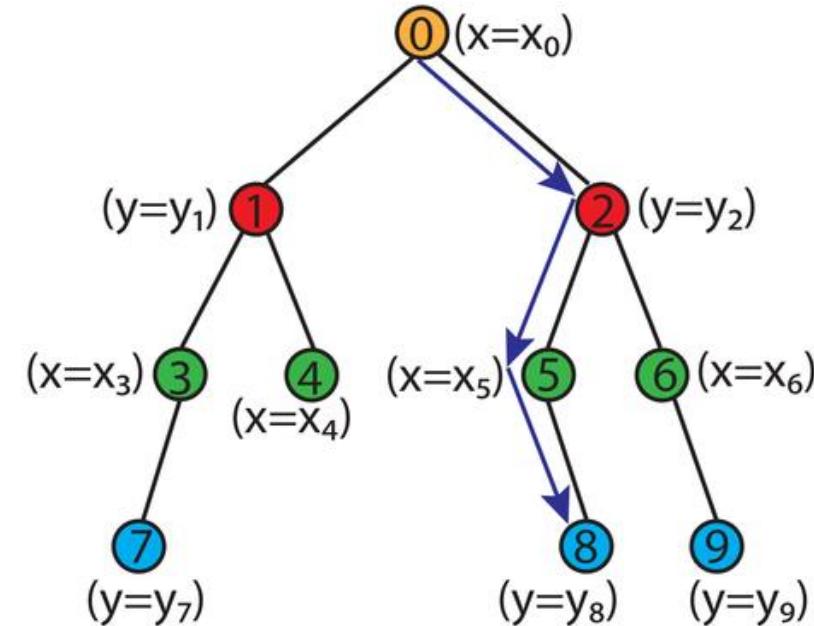


<http://wikipedia.org>

# Query in kd-trees



<http://www.mdpi.com/>



# kd-tree

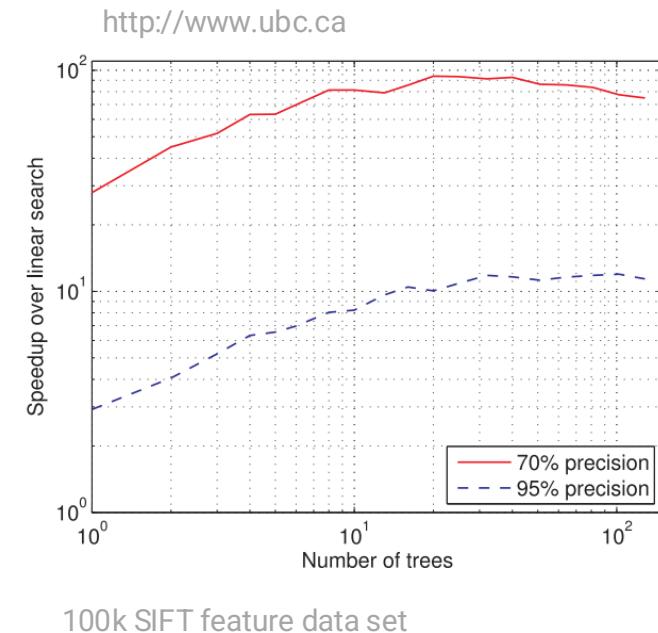
- Very **effective in low dimensional spaces**
- Performance **decreases for high dimensional data**
  - requires searching a large number of nodes
- Speedup by approximating queries:
  - ‘**error bound**’ approximate search
  - ‘**time bound**’ approximate search

# Randomized kd-Trees

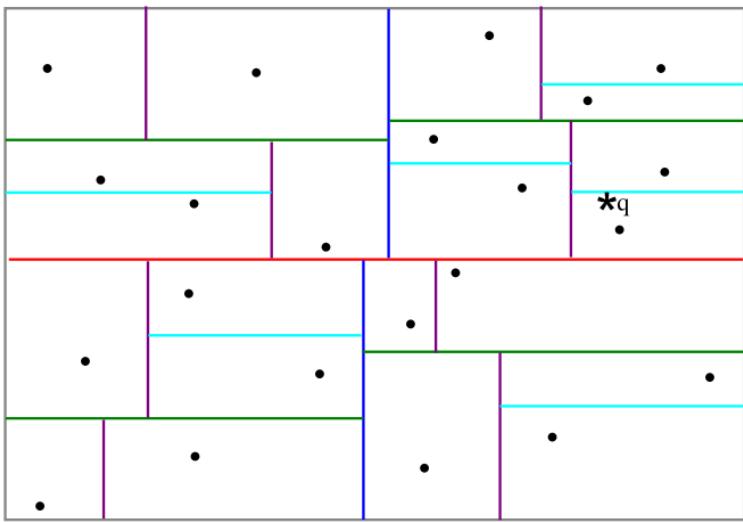
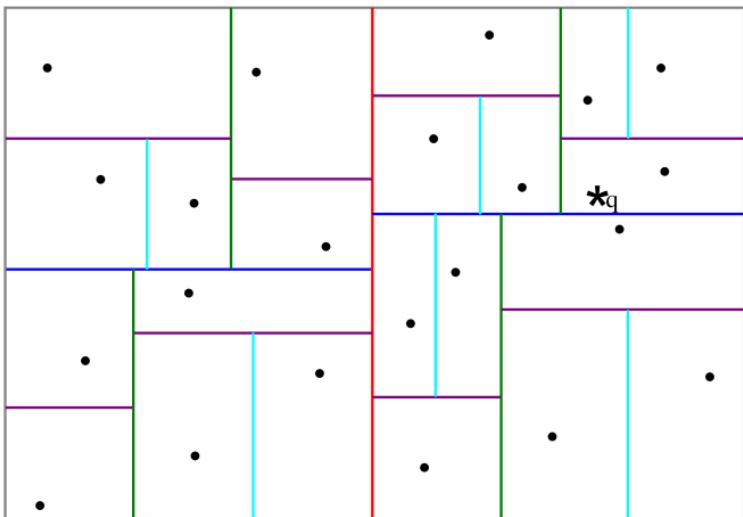
- **Idea:** multiple randomized k-d trees which are searched in parallel
  - build scheme similar to classic kd-tree
- **Difference to classic kd-tree:**
  - classic – split dimensions are chosen by highest variance
  - randomized – split dimensions are chosen randomly from the **top 5 dimensions highest variance**
- In practice often **20 trees** used

# Searching randomized k-d forest

- Maintain **single priority queue across forest**
- PQ is ordered by increasing distance
  - the search will explore first the closest leaves from all the trees
- examined **points are marked as visited across all trees** (not visited in another tree again)
- approximation is determined by **maximum number of visited leaf nodes**



# Randomized kd-Trees

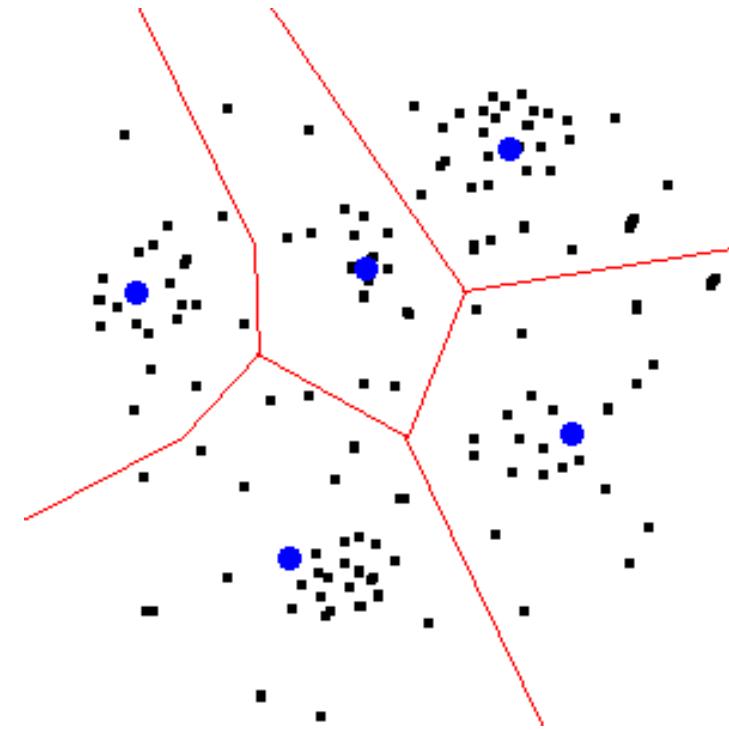


# **Priority Search k-means Tree**

# k-means

- important **clustering method** in CS
- partitions  $N$  data points in  $k$  clusters  $S = \{S_1, S_2, \dots, S_k\}$
- each data point belongs to cluster with nearest mean
  - mean of the cluster is cluster center/representative

$$\arg \min_s \sum_{i=1}^k \sum_{x \in s_i} \|x - \mu_i\|^2$$



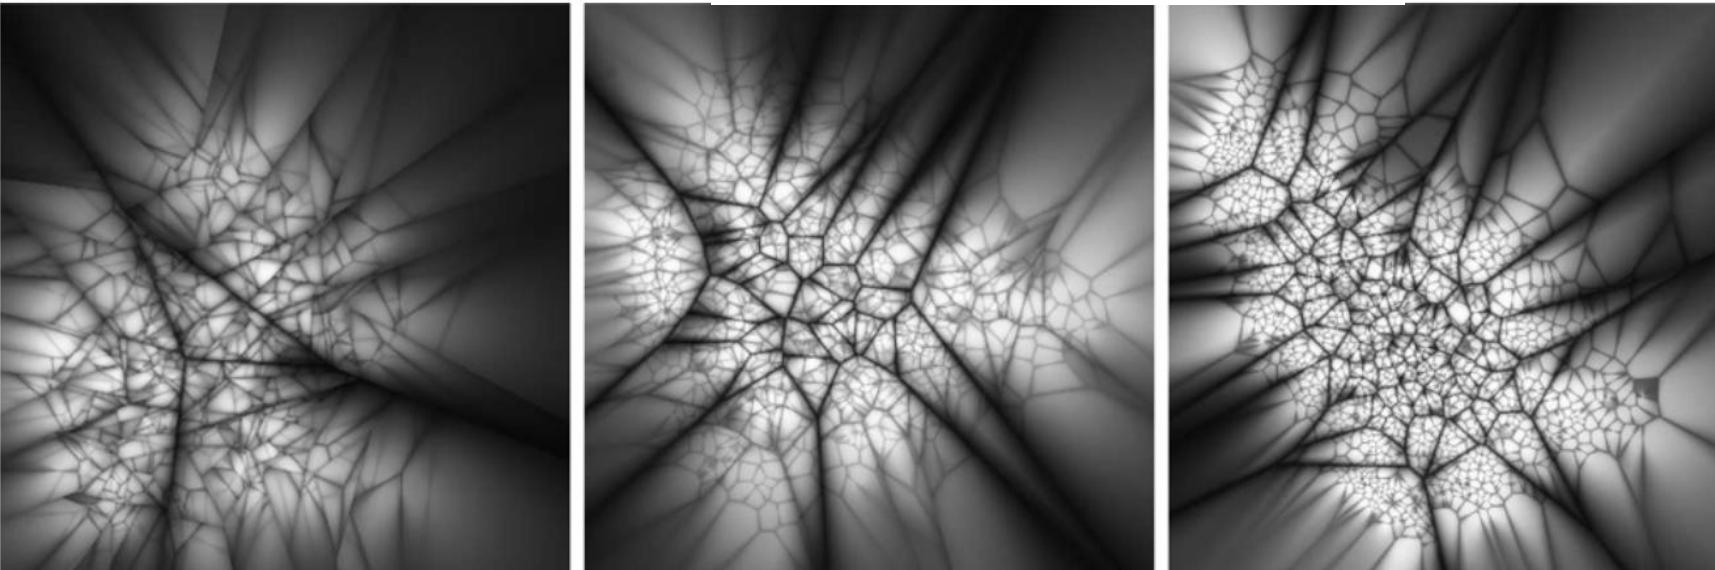
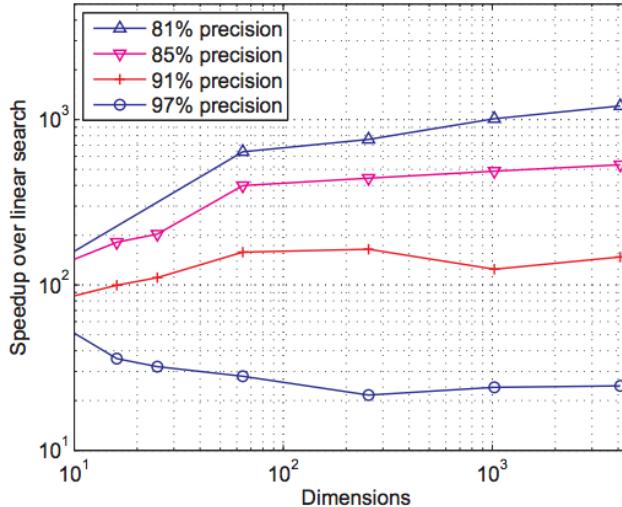
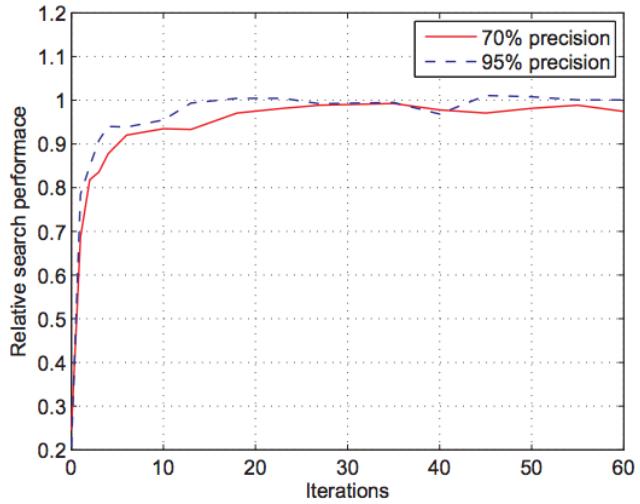
<http://mnemstudio.org/>

- results in partitioning of data space (Voronoi cells)

# k-means Trees

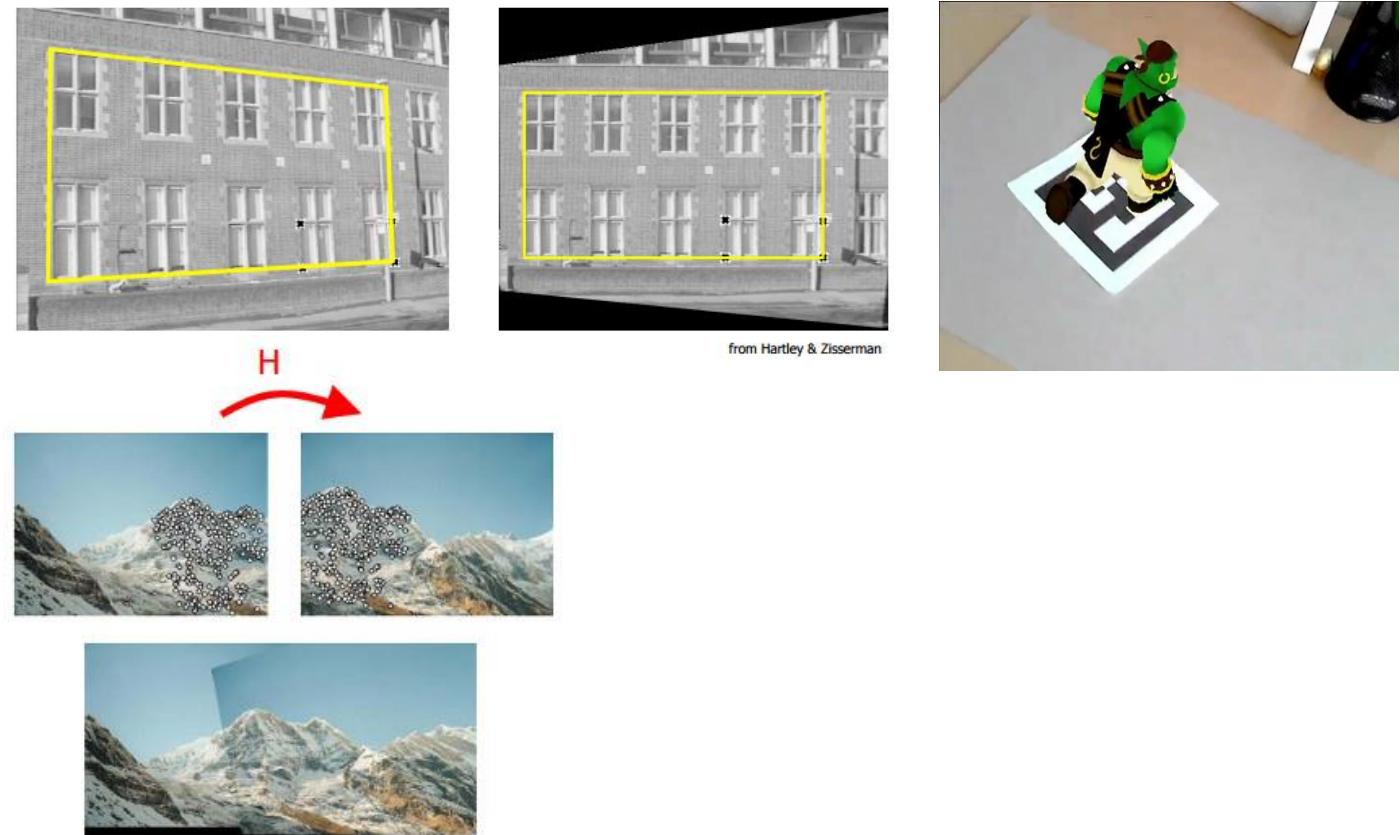
- **Construction:**
  - Tree is constructed by **partitioning the data points** at each **level into  $K$  distinct regions** using k-means clustering
    - Apply recursively to the points in each region
  - Recursion is stopped when the number of points in a region is smaller than  $K$
- **Search:**
  - initially traversing the tree from the root to the closest leaf (following closest cluster centre)
  - add unvisited branches to priority queue
  - use branch that has the closest center to the query point
  - restart the tree traversal from that branch
  - stop after number of visited leaf nodes

# k-means Trees

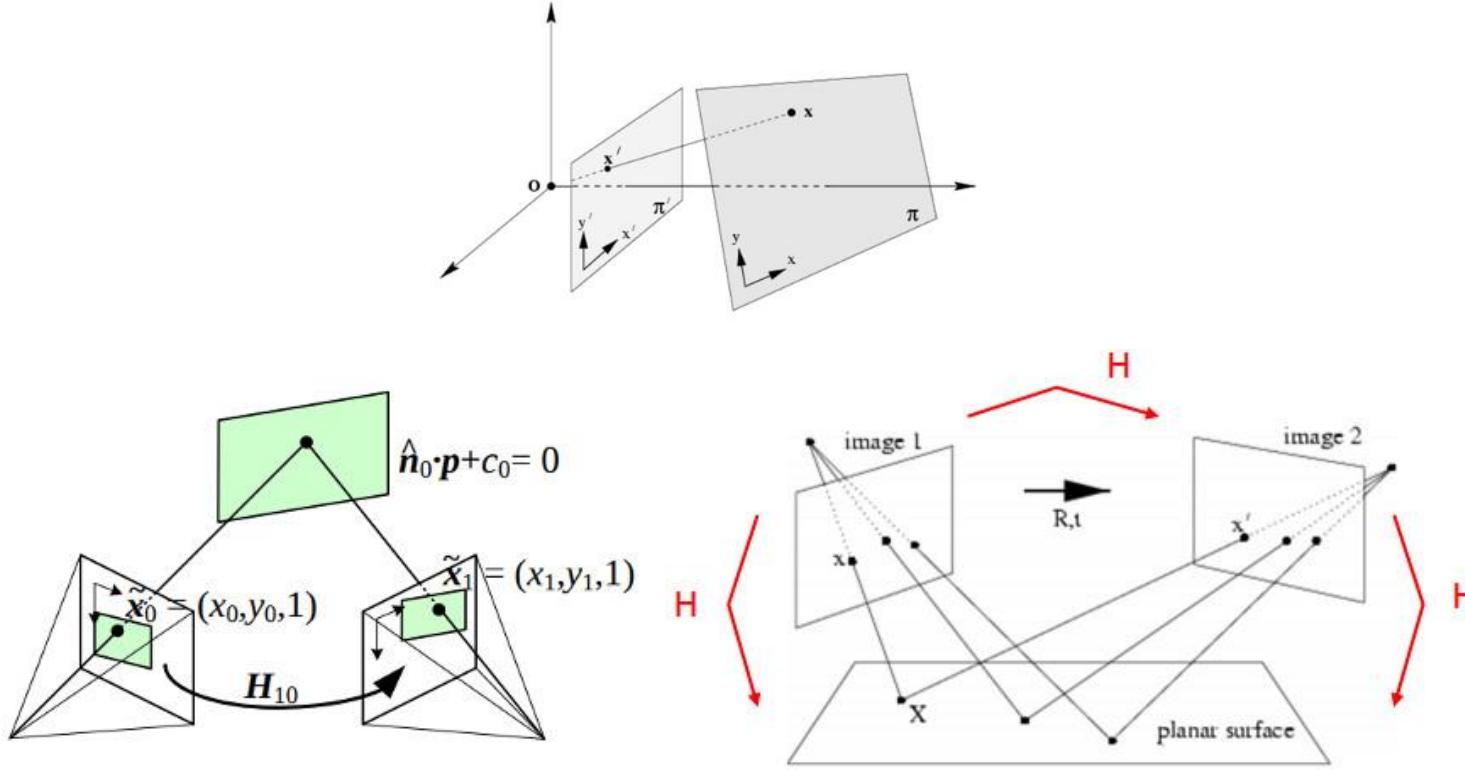


# Feature Matching Application: Homography

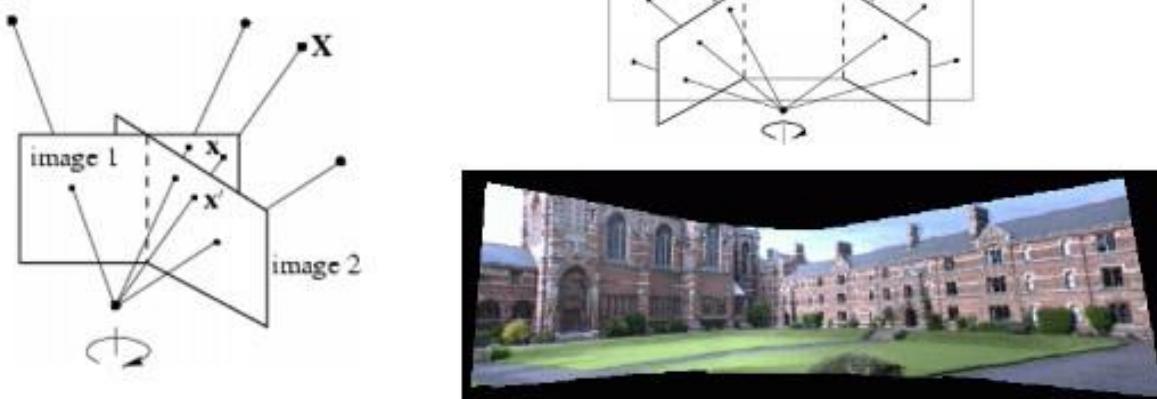
- Camera calibration
- 3D reconstruction
- Stereo vision
- Scene understanding



# Homography



Rotating camera, arbitrary world



# Homography matrix

$$\begin{bmatrix} {}^w x'_i \\ {}^w y'_i \\ w \end{bmatrix} \cong \begin{bmatrix} h_{00} & h_{01} & h_{02} \\ h_{10} & h_{11} & h_{12} \\ h_{20} & h_{21} & h_{22} \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix}$$

$$x'_i = \frac{h_{00}x_i + h_{01}y_i + h_{02}}{h_{20}x_i + h_{21}y_i + h_{22}}$$

$$y'_i = \frac{h_{10}x_i + h_{11}y_i + h_{12}}{h_{20}x_i + h_{21}y_i + h_{22}}$$

$$x'_i(h_{20}x_i + h_{21}y_i + h_{22}) = h_{00}x_i + h_{01}y_i + h_{02}$$

$$y'_i(h_{20}x_i + h_{21}y_i + h_{22}) = h_{10}x_i + h_{11}y_i + h_{12}$$

# Homography matrix

Lets do this in more detail  
and code

$$\begin{bmatrix} x_1 & y_1 & 1 & 0 & 0 & 0 & -x'_1 x_1 & -x'_1 y_1 & -x'_1 \\ 0 & 0 & 0 & x_1 & y_1 & 1 & -y'_1 x_1 & -y'_1 y_1 & -y'_1 \\ & & & & & : & & & \\ x_n & y_n & 1 & 0 & 0 & 0 & -x'_n x_n & -x'_n y_n & -x'_n \\ 0 & 0 & 0 & x_n & y_n & 1 & -y'_n x_n & -y'_n y_n & -y'_n \end{bmatrix} \begin{bmatrix} h_{00} \\ h_{01} \\ h_{02} \\ h_{10} \\ h_{11} \\ h_{12} \\ h_{20} \\ h_{21} \\ h_{22} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

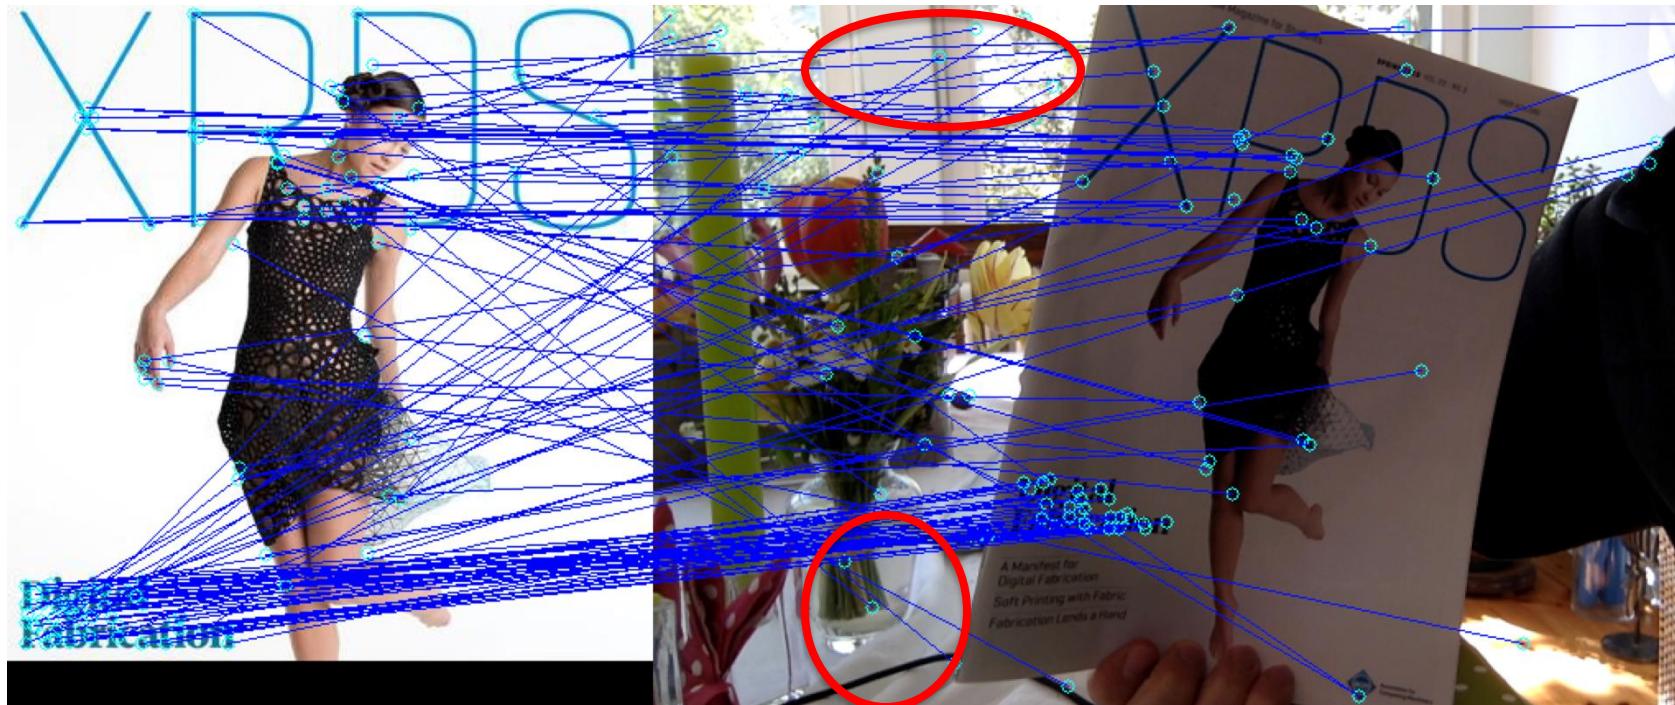
**A**  
 $2N \times 9$

**h**      **0**  
9       $2N$

- Defines a least squares problem:  $\min \|Ah - 0\|^2$ 
  - Solution:  $\hat{h}$  = eigenvector of  $A^T A$  with smallest eigenvalue
  - Works with 4 or more points

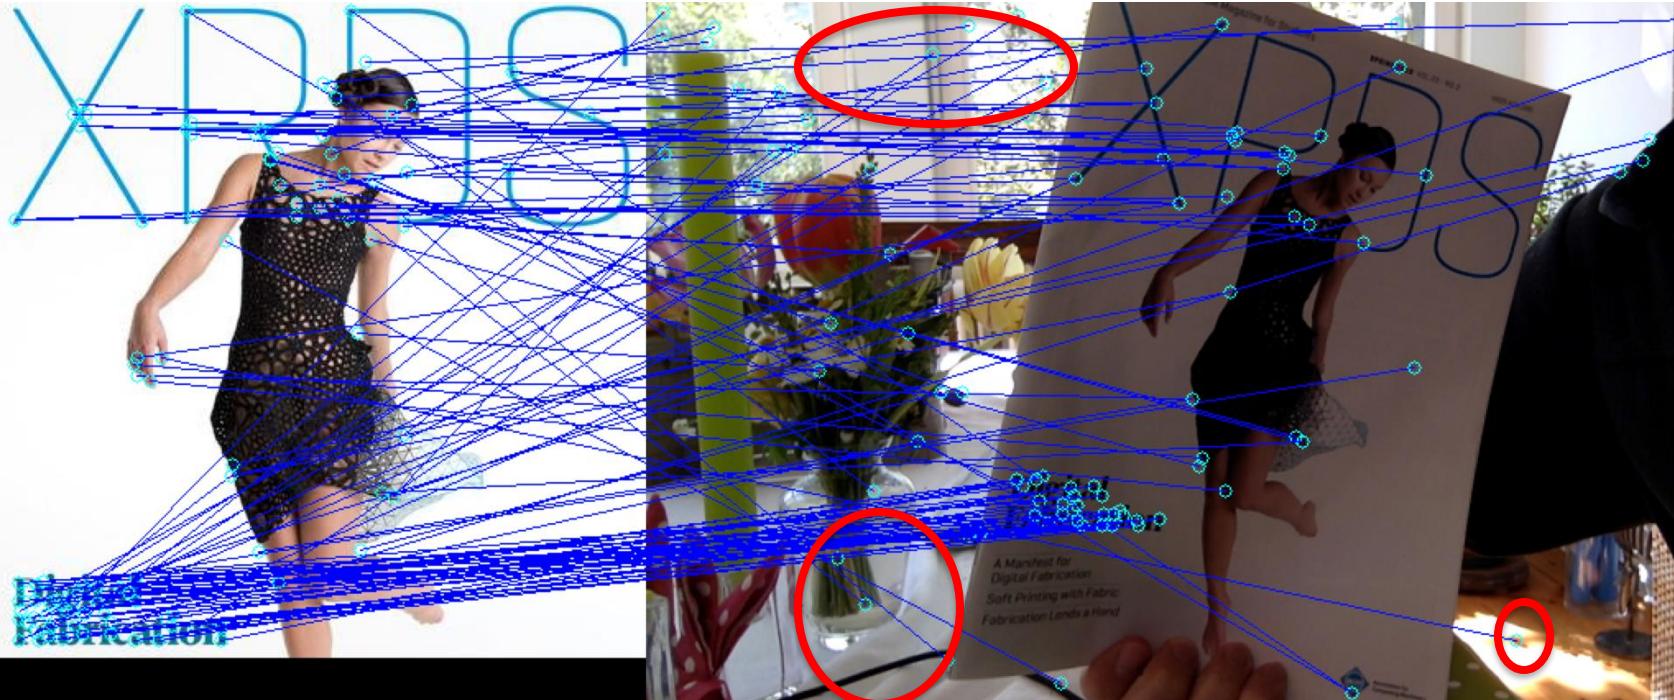
```
H, status = cv2.findHomography(pts_prevFrame, pts_curFrame, cv2.RANSAC, 3.0)
quad = cv2.perspectiveTransform(quad.reshape(1, -1, 2), H).reshape(-1, 2)
```

# RANSAC Motivation



# RANSAC

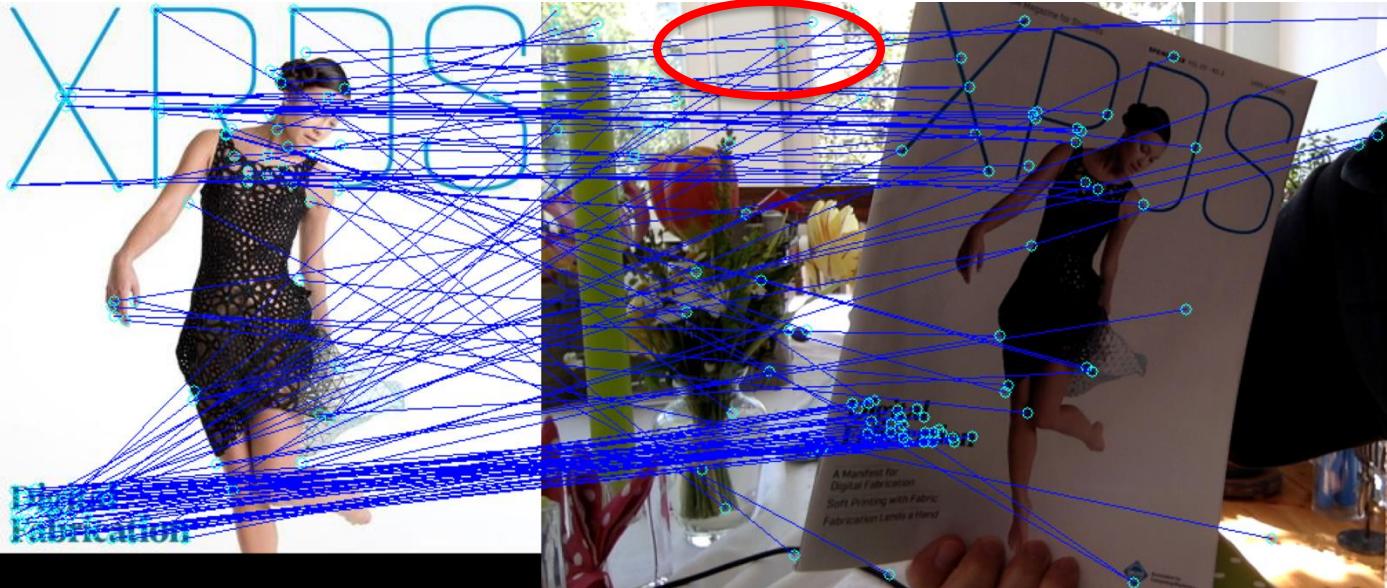
- What do we do about bad matches?



- Note: at this point we don't know which ones are good/bad

# RANSAC

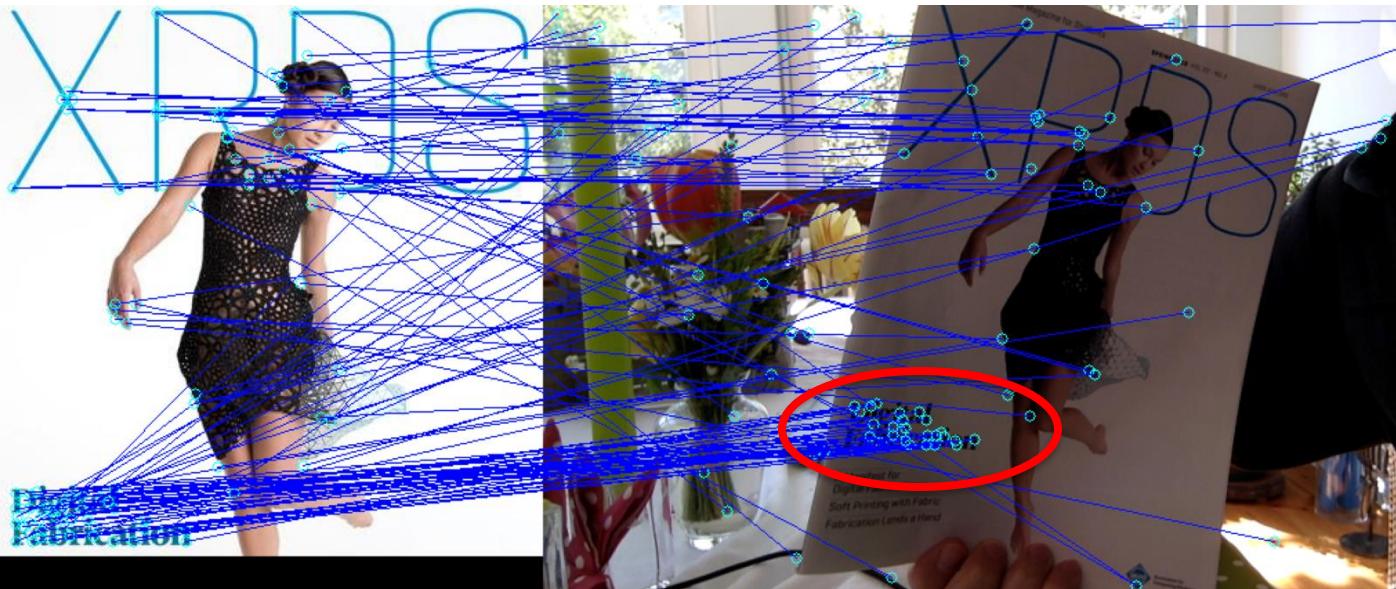
- Idea:
  - Select one match of at least 4 points and compute homography



- Count inliers using the homography matrix

# RANSAC

- Idea:
  - Select one match and count inliers



many inliers

- **Keep match with largest set of inliers**
- Find “average” transformation (homography  $H$ ), but with only inliers